

NTNU - NORGES TEKNISK-NATURVITENSKAPELIGE UNIVERSITET
Faculty of Engineering Science and Technology
Department of Civil and Transport Engineering
TBA4925 - Master Thesis

Optimizing the micro-tasking method and exploring its usage potential within geospatial data

Anne Sofie Strand Erichsen
Trondheim, June 2017

DAIM page

Background

Micro-tasking is an Internet phenomenon that has increased in popularity over the last years. The method is used for solving large tasks that can be divided into many smaller ones (micro-tasks). This involves both the use of computers and a large number of people. Importing geospatial data over large areas can be considered as a large task. This is a task that can be divided into smaller ones, for example by dividing into smaller areas.

Task Description

This Master thesis will have an emphasis on the data validation and conflict handling part of the import of geospatial data. These processes are too complicated to do fully automatic through scripts, and the thesis investigates how micro-tasking can be a relevant approach to the problem. Specific tasks:

- Study related literature
- Make a web-based experiment in order to answer the following questions:
 - What are the number of objects optimal within a micro-task to get it completed as quickly as possible?
 - Does the quality of the work vary between the different tasks given?
 - Do amateurs manage to do the tasks?
- Explore the micro-tasking methods usage potential within geospatial data

Administrative/guidance

The work on the Master Thesis starts on January 18th, 2017

The thesis report as described above shall be submitted digitally in DAIM at the latest at June 18th, 2017

External supervisor Atle Frenvik Sveen, Norkart AS

Supervisors at NTNU and professor in charge: Terje Midtbø

Abstract

Micro-tasking is the process of dividing a large task into smaller parts, making the task easier to distribute and lowering its complexity. This thesis proposes to extend the usage of the micro-tasking method to involve geospatial tasks. To the authors best knowledge little research has been conducted on how to best exploit micro-tasking together with geospatial data. This thesis explores how to partition a large task into smaller parts when it involves geospatial data and also analyses if individuals with no experience manage to solve these tasks.

An online web experiment was developed and implemented to gather data about how individuals solve geospatial micro-tasks. Statistical analysis is conducted on the collected data to answer the research questions. The author concludes that geospatial micro-tasks can be solved by inexperienced individuals and if the quality of the task results is important, fewer elements will increase the quality.

Sammendrag

Prosessen der en deler opp en større oppgave inn i mindre deler for å letttere distribuere oppgavene og redusere kompleksiteten kalles mikro-oppgaver. Forfatteren mener at mikro-oppgaver har et potensiale rundt geografiske oppgaver som idag ikke blir utnyttet. Det er lite, om noe, forskning på hvordan slike oppgaver lar seg løse gjennom denne oppdelings prosessen.

Masteroppgaven undersøker om mikro-oppgaver, som omhandler geografisk data, kan løses av alle mennesker, uavhengig av bakgrunnen deres og om antall elementer i hver mikro-oppgave påvirker kvaliteten på resultatet. Et eksperiment ble utviklet for å undersøke hvordan personer, både eksperter og uefarne, løser geografiske oppgaver delt inn i mikro-oppgaver. Den innhente dataen ble brukt i statistiske analyser for å teste ulike hypoteser. Masteroppgaven konkluderer med at det er mulig å bruke mikro-oppgaver i store geografiske oppgaver, og så lenge man har en introduksjonsdel som viser hvordan mikro-oppgavene skal løses kan alle, uavhengig av bakgrunn og utdannelse, løse disse.

Preface

This paper is a master thesis written for the Department of Civil and Transport Engineering at the Norwegian University of Science and Technology (NTNU) in Trondheim, Norway. It is a part of the study program Engineering and ICT - Geomatics and was written in the spring of 2017.

I would like to thank my supervisor Terje Midtbø for his help and feedback, and also Atle Frenvik Sveen for his support whenever I needed it. I would also like to thank my classmates for their invaluable feedback.

Trondheim, June 2017
Anne Sofie Strand Erichsen

Contents

Abstract	v
Sammendrag	vii
Preface	ix
1 Introduction	1
2 Background	3
2.1 Why do we need humans?	3
2.2 Human computation	4
2.3 Crowdsourcing	5
2.4 Micro-tasking	5
2.4.1 Micro-tasking platforms	6
2.4.2 Micro-tasking workforce	8
2.4.3 Micro-tasking usage	8
2.4.4 Building imports in OpenStreetMap using micro-tasking	10
2.4.5 Task complexity	11
2.4.6 Micro-tasking limitations	12
3 Experiment and Web application	15
3.1 Experiment	15
3.1.1 The two micro-tasks	15
3.1.2 The three experiment tasks	16
3.1.3 Building shapes	18
3.2 Web application	19
3.2.1 React application	19
3.2.2 Data acquisition	20
3.3 Pilot test	22
3.3.1 Execution of the pilot test	22
3.3.2 Results from the pilot test	23
3.4 Sample Size	24
4 Results	25
4.1 Participants	25
4.2 Statistics theory	26
4.2.1 Normal testing	26
4.2.2 Hypothesis testing	28
4.3 Survey results	32
4.3.1 Gathered data	32
4.3.2 Normality tests	35
4.3.3 Levene's test of Equality of Variance	41
4.3.4 Hypothesis testing	42

CONTENTS

5	Discussion	57
6	Conclusion	61
	Appendices	63
A	Normal tests	65
A.1	All participants, number of correctly chosen elements variable	65
A.2	All participants divided by task, total time variable	66
A.3	All participants divided by task, correct element variable	67
A.4	Experienced participants divided by task, total time variable	68
A.5	Experienced participants divided by task, correct elements variable	70
A.6	Inexperienced participants divided by task, total time variable.	70
A.7	Inexperienced participants divided by task, correct elements variable	71

List of Figures

2.1	Illustration of how a large task can be partitioned and distributed through micro-tasking (Michelucci and Dickinson, 2016)	6
2.2	Micro-tasking workers task selection process (Schulze et al., 2012)	8
2.3	Crisis map (Meier, 2014)	9
2.4	Swimming pools (Nikki, 2016)	10
3.1	Question one as it is displayed in the web application during task B	16
3.2	Question two as it is displayed in the web application during task B	16
3.3	Illustration of the three tasks given in the experiment	17
3.4	Creating of footprint layers used in question one	18
3.5	Interface for the client application	20
3.6	Population vs. sample	24
4.1	The distribution of the task order in the analysed data	26
4.2	Skew (MedCalc Software bvba, 2017)	27
4.3	Kurtois (MedCalc Software bvba, 2017)	27
4.4	Two-sample t-test hypotheses	29
4.5	One-way ANOVA hypotheses	30
4.6	Histograms with normal distribution fit with samples containing total time to complete each task	36
4.7	Histograms with normal distribution fit after Box-Cox power transformation	37
4.8	Histogram with normal distribution fit	38
4.9	Histogram with normal distribution fit after Box-Cox	39
4.10	Sample 1 and 2 - mean (green dot) and standard deviation (blue line)	44
4.11	Sample 3 and 4 - mean (green dot) and standard deviation (blue line)	46
4.12	Sample 5, 6 and 7 - mean (green dot) and standard deviation (blue line)	48
4.13	Sample 8, 9 and 10 - mean (green dot) and standard deviation (blue line)	50
4.14	Sample 11, 12 and 13 - mean (green dot) and standard deviation (blue line)	51
4.15	Sample 14, 15 and 16 - mean (green dot) and standard deviation (blue line)	52
4.16	Mean (green dot) and standard deviation (blue line) for sample 17, 18 and 19	53
4.17	Mean (green dot) and standard deviation (blue line) for sample 20, 21 and 22	54
A.1	Histograms with normal distribution fit with samples containing the number of correctly chosen elements	65
A.2	Histogram with normal distribution fit - sample with total time per task	66
A.3	Histogram with normal distribution fit after Box-Cox transformation, total time variable	67

LIST OF FIGURES

A.4 Histogram with normal distribution fit showing samples with number of correct elements per task	67
A.5 Histogram with normal distribution fit after Box-Cox	68
A.6 Histograms with normal distribution fit with samples containing total time to complete each task	69
A.7 Histograms with normal distribution fit containing Box-Cox transformed data	69
A.8 Histogram with normal distribution fit showing samples with number of correct elements results for experienced participants	70
A.9 Histogram with normal distribution fit	71
A.10 Histogram with normal distribution fit after Box-Cox	71
A.11 Histogram with normal distribution fit	72
A.12 Histogram with normal distribution fit after Box-Cox transformation	72

List of Tables

3.1	Experiments variables	21
4.1	Total time, all participants	32
4.2	Correct elements, all participants	33
4.3	Total time, divided by task	33
4.4	Correct elements, divided by task	33
4.5	Total time, divided by task, only experienced	34
4.6	Correct elements, divided by task, only inexperienced	34
4.7	Total time, inexperienced per task	35
4.8	Correct elements, inexperienced per task	35
4.9	Summary, normality tests	40
4.10	Summary, Levene's tests	41
4.11	Summary, hypothesis tests	55

1 | Introduction

Having access to information is not the same as learning and understanding it. Today, the amount of information available is enormous. Computers have the ability to learn through machine-learning but are often dependent on well-developed training data to learn. The data has to be information placed in contextual meaning. Understanding and giving contextual meaning to information is the strengths of the human brain. Humans have the ability to create new ideas and concepts from unstructured information (Ross and Jamilly, 2016). Computers and the human brain are an unbeatable combination. Humans need computers for their speed and accuracy, and computers need the human brain to make sense of new information.

Albert Einstein illustrates this perfectly: “Computers are incredibly fast, accurate, but stupid. Humans are incredibly slow, inaccurate, but brilliant. Together they may be powerful beyond imagination” (Holzinger, 2013).

This thesis argues that a useful approach for combining the strengths of computers and humans together is micro-tasking. Micro-tasks are the smallest, simplest types of tasks and should demand little time to complete. The tasks can be carried out by humans through, or in collaboration with, computer systems (Yang et al., 2016). The author believes that micro-tasking has an unexplored potential within geospatial tasks, that is, tasks that involve data with a geographic position. Importing geospatial data into a database, i.e., OpenStreetMap, covering huge areas can be considered a large task. This thesis will have an emphasis on the data validation and conflict handling part of the import of geospatial data.

Based on careful reading of relevant literature, the author conclude that little research has been done on micro-tasks involving geospatial data. To be able to exploit the micro-tasking method together with geospatial data, it is important to study how well humans solves these kinds of tasks. Micro-tasks are often published on a micro-tasking platform, where tasks and humans are connected. It is important to know if inexperienced individuals are capable of solving the tasks. If only experienced people can solve geospatial micro-tasks, a platform who can distinguish between people with geospatial knowledge needs to be used. The author does not have any knowledge if such a platform exists.

In Remote-Sensing, humans perform land-cover classification tasks (Salk et al., 2016). At least two papers have studied whether there are significant differences in quality of the information contributed by experts and non-experts. Salk et al. (2016) concluded that there was little first-order relationship between professional background and the task accuracy. When comparing specialists with non-specialists, the non-specialists performed slightly better on images near home. See et al. (2013) concluded that there was little difference between experts and non-experts, and also found that the non-experts improved more than experts over time. Overall the non-experts were as reliable in what they identified as the experts. These two studies show that the person’s background is not necessarily important. See et al. (2013) argued that with

1. INTRODUCTION

proper targeted training material the differences between experts and non-experts could potentially decrease.

The goal of this thesis is to study if micro-tasking can successfully be expanded to involving maps and geospatial data. The OpenStreetMap (OSM) community has used the method for some years, and the usage so far can be evaluated as successful (Erichsen, 2016). This thesis will look at the OSM community's usage of micro-tasking to examine if all individuals, independent of background, manage to solve geospatial tasks. The thesis also aims to determine if the number of elements in each micro-task has an impact on time spent per task and the quality of the task results.

There is potentially much work creating micro-tasks. The large task needs to be appropriately broken down to micro-tasks that are easy, enjoyable, and fast to solve. This breakdown requires design skills and proper tutorials and examples for new workers (Schulze et al., 2012). Guidelines can be used to avoid putting too much emphasis on the preparations. This thesis can give a set of guidelines to be utilized on how to break down large geospatial tasks.

The aim of this thesis is to answer the research questions:

1. Is it possible to give micro-tasks containing geospatial data to inexperienced workers?
2. Will the quality of the solved task increase with fewer elements present in each micro-task?
3. What is the number of elements optimal within a micro-task to get it completed as quickly as possible?

An online web experiment was developed to answer the research questions. The experiment contains three tasks varying the number of elements given to the participant. Each task has the same two questions representing two micro-tasks. One question involves map interaction and the other an interpretation task containing metadata.

The next chapter will give a thorough introduction to micro-tasking and clarify the aim of this thesis. Chapter three will introduce the experiment, the web application hosting the experiment and results from the experiments pilot test. Chapter four will present the individuals participating in the experiment, provide a description of the statistical theory used in the analysis, and give a summary of the gathered data. The last section in chapter four will contain the analyses of hypotheses which in sum will answer the research questions. Chapter five contains the discussion, where the results are summed and discussed. At last, the thesis will give a conclusion and outline further work in this area.

2 | Background

Creating and maintaining real-world knowledge bases in a classical work environment demands a high cost, and is a cost that is often unnecessary (Meier, 2013b, p. 134). Alternative approaches are to rely on the knowledge of open crowds, volunteer contributions, or services like micro-tasking platforms where there are people ready to work on the tasks given to them (Meier, 2013b, p. 134).

Today, geospatial data is more available than ever. Governments are releasing more and more data, and the OpenStreetMap database is still growing. While general data availability is increasing, the quality of the data is not necessarily perfect, and manual pre-processing is often necessary before use (Difallah et al., 2015). Pre-processing of the data is time-consuming and expensive. By exploiting both machines and people through the appropriate platform and approach, the cost can decrease and the quality increase. This thesis will argue that combining machines and people is often a better and faster solution than a fully-automatic or fully-manual approach. Implementing this kind of approach into a micro-tasking platform can be a good solution.

2.1 Why do we need humans?

Machine learning gives computers the ability to learn without being explicitly programmed. It involves computer intelligence, but the computers do not know the answers up front (Stanford University, 2017). Machine-learning algorithms have enormous problems when contextual information is missing. Without a pre-set of rules, a machine has trouble solving the problem. Machines do not have creativity, which is required to answer complex problems (Holzinger et al., 2016). According to the company Mighty AI¹, humans cannot be removed from Artificial Intelligence training loops. Machine-learning approaches still require a huge amount of training data to work on new domains (Schulze et al., 2012). Mighty AI believe that humans will continue to play a crucial role in creating training data for the algorithms (Gutzwiller, 2017).

It is suggested by Biewald (2015) and Oppenheimer (2017) that machine learning accuracy should follow the Pareto 80:20 principle. Getting 80 % accuracy can be reasonably easy to accomplish, but the last 20 % should be handled by human input (Biewald, 2015; Oppenheimer, 2017). Significant human input is to provide training data and route tasks when the algorithm is unsure of its answer (Nakhuda, 2016; Oppenheimer, 2017). The machine learning company developmentSEED² use a micro-tasking solution to clean their machine learning output data. They are using humans to get faster and more accurate output data by developing Skynet Scrubber, a GUI

¹Mighty AI generates high-quality AI training data.

²DevelopmentSEED is a creative engineering team solving complex problems with open software and open data.

2. BACKGROUND

web application solution to get human input quicker and easier³. In their blog, Derek Lieu writes: "Skynet gets more capable every day, but the output is still not perfect [...] We built Skynet Scrub so we could start using Skynet data sooner" (Lieu, 2017).

Holzinger (2016) claim that most people from the machine learning community are concentrating on *automatic* machine learning by bringing the humans out of the process. When humans are out of the loop, the training data sets can be uncertain and incomplete, and the resulting algorithm can be questionable (Holzinger, 2016). By bringing humans back into the process, especially in domains where the data sets are questionable, for instance in the health domain, one enables what neither a human or a computer can do on their own (Holzinger, 2016). It is possible to build hybrid human-machine systems that combine both the scalability of computers and the yet unmatched cognitive abilities of the human brain (Difallah et al., 2016). The co-founder of Palantir Technologies⁴ stated that "Computers are bad at finding patterns unless we have a well-understood problem" (Cohen, 2013). Only humans can understand and frame a new problem. Palantir Technologies believe in augmenting human intelligence, not replacing it. As Holzinger (2013) say, "[...] the problem-solving knowledge is located in the human mind and - not in machines, " and this is something we must acknowledge.

2.2 Human computation

Human computing is, at its most general level, computation performed by human beings and a human computation system contains both humans and computers working together to solve difficult problems (Schulze et al., 2012). We argue that utilizing the human processing power is still important. Humans are necessary even though our computers are becoming more and more complex. Traditional approaches are to focus on improving the software when solving computational difficulties, but as the reader will see in this thesis, a solution that uses humans cleverly by exploiting the cognitive abilities of the human brain can create much faster and better results than software. One of the pioneers of crowdsourcing, Luis von Ahn, wanted to find a cheap and efficient way to label images (von Ahn, 2008). The solution was to exploit the use of a game-like approach in a non-game context to motivate individuals to label the pictures through a game. This method is called gamification (Huotari and Hamari, 2017). The game was called "The ESP game" and solved the problem of labeling images with words. Most images do not have a proper caption associated with them, and this makes it difficult to create search engines for images. A fast and cheap method of labeling images is by using humans cleverly, and humans can very easily see if the picture contains, i.e., a dog or a cat. Through "The ESP game" humans were labeling images without even knowing it, they only played a fun game. Within a few months, the game collected more than 40 million image labels (von Ahn, 2008), and they did not even have to pay them to do it. Human computation is one of the major areas

³developmentSEED's algorithm is called Skynet

⁴Palantir builds software that connects data, technologies, humans and environments

where the gamification approach has been employed (Morschheuser et al., 2016). Each human performs a small part of a massive computation task.

Human computation, a term introduced by Luis von Ahn, refers to, according to Quinn and Bederson (2011), a distributed system that combine the strengths of humans and computers to accomplish tasks that neither can do alone. To make human computation through crowdsourcing compelling, one needs to know how the results are optimally acquired from humans and how to integrated the results into productive environments without having to change established workflows and practices (Meier, 2013b, p. 134). Gamification can be one solution on how to make human computation in crowdsourcing effective (Wang et al., 2017). The author will argue that micro-tasking can be another solution for effective crowdsourcing using the cognitive abilities of humans.

2.3 Crowdsourcing

The first time the term "crowdsourcing" appeared was in a Wired magazine article by Jeff Howe (Howe, 2006). Whereas human computing (section 2.2) replaces computers with humans, crowdsourcing replaces traditional human workers with members of the public (Quinn and Bederson, 2011). Crowdsourcing companies came onto the scene with the goal of helping businesses solve simple problems on a massive scale (Webster, 2016). EYeka (2015) state that 85 % of the top global brands use crowdsourcing for various purposes. Crowdsourcing is an increasingly important concept (Salk et al., 2016) and has become a widespread approach to dealing with machine-based computations where we leverage the human intelligence (Gadiraju et al., 2015a). Crowdsourcing is a way of refactoring work in a manner that exploits the worker's flexibility. It focuses on getting the right skills to the right part of the problem. Dividing a large task into smaller parts makes it easier to distribute the smaller tasks to the right person and expertise. The partitioning and distribution can be accomplished through micro-tasking, also called "smart crowdsourcing" by Patrick Meier (Meier, 2013a).

When the field of a crowdsourced project is explicitly geographical, it is often called *volunteered geographical information* (VGI). According to Salk et al. (2016), the best known VGI project is OpenStreetMap (OSM). OSM is an open-source mapping project, where volunteers contribute with their local knowledge and mapping abilities, called "the Wikipedia of maps" (Palen et al., 2015). Wikipedia can be said to be the best-known and most successful example of crowdsourcing on a global scale. Chilton et al. (2009) suggested that the OSM project would have a similar impact on open and free geospatial data as Wikipedia had on 'fact finding'.

2.4 Micro-tasking

The simplest type of tasks are called micro-tasks and is illustrated in figure 2.1. Micro-tasks should not require any special training, and a task should be completed within a couple of minutes (Ipeirotis and G., 2010). Problems that are suitable for solv-

2. BACKGROUND

ing through micro-tasking are those that are easy to distribute into many simple tasks, which can be completed in parallel in a relatively short period of time, and that does not require specific skills (Sarasua et al., 2012). Research has also demonstrated that micro-tasking is effective for far more complex problems when using sophisticated workflow management techniques. Micro-tasking can then be applied to a broader range of challenges like: 1) completing surveys, 2) translating text between two languages, 3) matching pictures of people, 4) summarizing text (Bernstein et al., 2015).

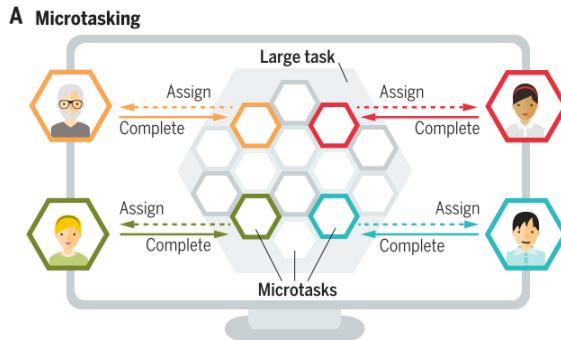


Figure 2.1: Illustration of how a large task can be partitioned and distributed through micro-tasking (Michelucci and Dickinson, 2016)

2.4.1 Micro-tasking platforms

Several micro-tasking platforms with large user-bases are found online. These platforms serve as a place where the crowd (workers) willing to perform small tasks is connected with work providers (Difallah et al., 2015). Tasks conducted by the crowd has to pass a quality control mechanism to be approved. A common approach to ensure reliability in the answers is by collecting multiple judgements from the crowd (Gadiraju et al., 2015b). The correct answer is then determined by the majorities answer. The quality mechanisms vary between the platforms, but a common measure, noted by the author, is that at least three independent workers need to agree on the task before approved. This was also observed by James McAndrew⁵. Three popular platforms are presented in this section.

2.4.1.1 Amazon's Mechanical Turk

Amazon Mechanical Turk (MTurk) is a very popular micro-tasking platform created in 2005 and is still in use today (Difallah et al., 2016). MTurk acts as an online labor marketplace (Sarasua et al., 2012). It provides the infrastructure, connectivity and

⁵James McAndrew noted it in a email correspondence with the author

payment mechanisms so that hundreds of thousands of people can perform micro-tasks on the Internet and get paid per completed task. MTurk is used for many different tasks that are easier for people than computers. It contains simple tasks such as labeling or segmenting images or tagging content, to more complex tasks such as translating or even editing text (Franklin et al., 2011). In the marketplace, employers are known as requesters, and they post tasks, called *human intelligence tasks* (HIT's). The HIT's are then picked up by online users, *crowd workers*, who complete the tasks in exchange for a small payment (a few cents per HIT) (Ipeirotis and G., 2010). The workers select the HIT's themselves, a fundamental difference from traditional employment models (Schulze et al., 2012).

2.4.1.2 Tasking manager

The Tasking Manager tool is OpenStreetMap's micro-tasking platform. It was created in the aftermath of the Haiti earthquake in 2010 (Palen et al., 2015). The tool is used to coordinate satellite image tracing projects and sort the area covered by the satellite image into grids so that multiple people can map the same area at the same time. Each person works at one grid each. This way they avoid mapping the same areas. Partitioning the areas into grids is a very effective approach to coordinate a mapping job. The tasking manager is mainly used by the *Humanitarian OpenStreetMap Team* (HOT). This platform does not have a rewarding system or a gamification approach. It is solely based on volunteer contributors. This platform shows that it is not necessary to have a game or reward for to be successful.

There are also other micro-tasking tools in OpenStreetMap. MapRoulette and To-Fix are examples of such tools. Both tools are listed as error detection tools on the OSM quality assurance wiki page. MapRoulette uses a gamification approach, while To-Fix does not. The tools break common errors in the data into micro-tasks so that multiple individuals can work on the tasks simultaneously.

2.4.1.3 CrowdFlower

CrowdFlower is a company that wants to help businesses take advantage of crowdsourcing and human computation. They act as an intermediary for these companies (Quinn and Bederson, 2011). CrowdFlower receives tasks from businesses wanting to crowdsource their work or problems. CrowdFlower operates with a variety of services to get connected with workers (i.e., MTurk) (Quinn and Bederson, 2011).

What is special with CorwdFlower is their close ties with AI technology and a crowd-sourced workforce. Their costumers are allowed to perform tasks with algorithms and machine learning, but also introduce human judgement when they are not confident in the technology, and the human work can make the algorithms smarter (Ha, 2016). The founder of CrowdFlower says that "self-driving cars have gotten pretty good at recognizing many of the objects they encounter on the street, [...] (but) they can still

2. BACKGROUND

struggle with tricky things like “a person in a Halloween costume dressed as a stationary object, or a pole with a person painted on it,” which is where CrowdFlower comes in.” (Ha, 2016). This is a good example of why humans judgement is necessary and needs to be involved.

2.4.2 Micro-tasking workforce

It is said that crowdsourcing is radically changing the nature of work (Deng et al., 2016). Traditional workers are restricted to offices and arranged office hours. With crowdsourcing, through for instance micro-tasking platforms, the workers can choose when to work, and even better: which jobs to perform. This appears very attractive, but is it only on the surface?

According to Deng et al. (2016), crowdsourcing is radically changing people’s perspectives on how to manage their work-life balance. Compared to ”traditional” work tasks, the micro-tasks are simple and fast to finish (within a couple of minutes). The worker is also often motivated by tiny rewards every time they complete a micro-task.

Individuals who perform micro-tasks for micropayment is called *crowd workers* by (Deng et al., 2016). A study done on workers in the micro-tasking platform MTurk (section 2.4.1.1), says that the workers are representative for the general Internet user population, but are generally younger and have lower incomes and smaller families (Ipeirotis and G., 2010). Workers select which tasks to solve themselves. A task selection process is visualised in figure 2.2. Micro-tasks can also be assigned to workers, as shown in figure 2.1. This can be beneficial if the task requires specific skills as education or language (Schulze et al., 2012). Most workers look for tasks that utilize their knowledge, skills, and abilities in the best possible way (Schulze et al., 2012).

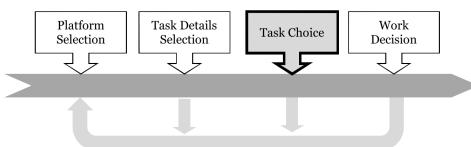


Figure 2.2: Micro-tasking workers task selection process (Schulze et al., 2012)

2.4.3 Micro-tasking usage

Micro-tasking and human computation have close ties. In the "Handbook of Human Computation", micro-tasking is strongly present in the *Human Computation for Disaster Response* chapter (Meier, 2013b, p. 95-105), as well as in several other parts of the book. In the disaster response chapter, the authors give an overview of how human computation methods, such as paid micro-tasks, could be used to help in major disasters. In 2012, the Philippines was hit by a typhoon called Ruby, devastating large regions. Through CrowdFlower the workers collected over 20 000 tweets related to the typhoon and identified the tweets containing links to either photos or video footage

from the damaged areas. The photos and videos in the relevant tweets were tagged and geo-tagged by volunteers if they portrayed evidence of damage. Within 12 hours a dataset of 100 georeferenced images and videos were collected. It resulted in a very detailed crisis map shown in figure 2.3. This map was the first official crisis-map based solely on social media content (Meier, 2013b, p. 101). In the aftermath of this crisis, an algorithm was developed to automatically detect tweets that link to photos and videos, which freed more time for the volunteers to georeference and tag more images and videos portraying evidence of damage (Meier, 2014).

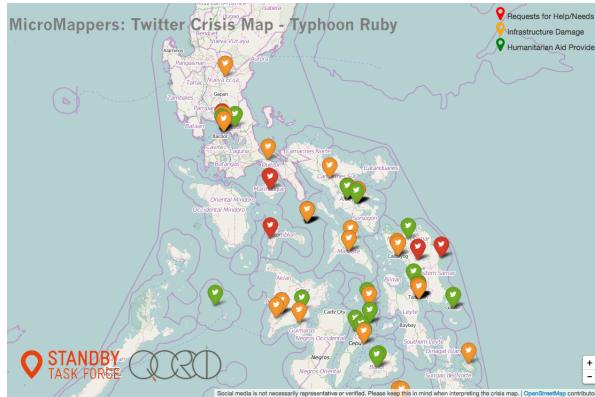


Figure 2.3: Typhoon Ruby Crisis map (Meier, 2014)

The typhoon Ruby crisis map is a good example of exploiting human computation, with crowdsourcing and micro-tasking. It refers to a problem-solving model where a problem or task is outsourced to a distributed group of people by splitting the task or problem into smaller sub-tasks or sub-problems. The sub-tasks are solved by multiple workers independently, often in return for a reward (Sarasua et al., 2012). Thanks to micro-tasking platforms as CrowdFlower and MTurk, it is possible to build a hybrid human-machine system that combines the scalability of computers with the yet unmatched cognitive abilities of the human brain (Difallah et al., 2016). When analyzing data from MTurk, findings from Gadiraju et al. (2015a) indicate rapid growth in micro-task crowdsourcing. With the establishment of micro-task crowdsourcing platforms as MTurk and CrowdFlower, micro-tasking is much more accessible. Micro-tasking practitioners are actively turning towards paid crowdsourcing to solve data-centric tasks that require human input (Gadiraju et al., 2015a). Most cases of micro-tasking combine human computation abilities with crowdsourcing.

Companies developing machine-learning algorithms have also seen an advantage with combining fast machine learning algorithms with human computation abilities and crowdsourcing. An example is the team Tomnod⁶, who did a project in Australia where they combined human computation, crowdsourcing, and machine learning to

⁶Tomnod is a team of volunteers who work together to identify important objects and interesting places in satellite images; www.tomnod.com

2. BACKGROUND

locate swimming pools (Kostas, 2016). The machine learning algorithm classified polygons where it was likely to be a swimming pool inside. The crowd who participated in finding the swimming pools were only presented with the classified polygons, which minimized the search area and then also the time required finishing the job (Kostas, 2016). Examples of classified polygons are shown in figure 2.4. The Tomnod team micro-tasked the work. The resulting dataset was then used to train a swimming pool detecting convolutional neural network (Nikki, 2016).



Figure 2.4: polygons containing swimming pools classified by the algorithm (Nikki, 2016)

Another example of micro-tasking usage is the New York Public Library. They use micro-tasking to train computers to recognize building shapes and other data on digitized insurance atlases. Humans complete micro-tasks where they check the computer's work and also capture information the computer missed. Individuals contributing checks and fixes building footprints drawn by the computer. The individuals also enter addresses and classify the building footprints using colors. To ensure accuracy, the same footprints are shown to several people. At least three different individuals check the same footprint, and 75% or more must agree on the footprint for the answer to be approved.

Most cases of micro-tasking usage exploit the large volume capabilities machines have and the cognitive capabilities of humans (Difallah et al., 2016). One of the advantages of micro-tasking platforms like MTurk, Tomnod and CrowdFlower, mentioned by (Meier, 2013b, p. 99), is the built-in quality control mechanisms that ensure a relatively high quality of output data. They set a review constraint, for instance in a project where they tagged satellite imagery of Somalia. Each unique image was reviewed by at least three different volunteers and only when all three agreed on type and location it was approved.

2.4.4 Building imports in OpenStreetMap using micro-tasking

In OpenStreetMap (OSM), at least two large building imports have been successfully achieved using micro-tasking (Erichsen, 2016). The first was an import in New York, the second in Los Angeles. Both import teams divided the building dataset into

smaller parts to lower the complexity. They used the same Python script to create the micro-tasks containing small chunks of building data. Having small chunks of building data made it possible to review and import the data manually into the OSM database (Barth, 2014a). In New York, they imported one million buildings partitioned into 5258 micro-tasks (Barth, 2014a). In Los Angeles, their dataset contained three million buildings (Sambale, 2016). We have not found how many micro-tasks that were created in total. In both projects, all buildings needed to be quality checked and merged correctly with existing data in OSM. This validation process is done manually in OSM (community, 2017a). Both building imports used the tasking manager platform (2.4.1.2) to organize and distribute the micro-tasks (Barth, 2014b) (community, 2017b). The number of buildings in each micro-task varied because the building dataset was partitioned based on already existing subregions which had varying building densities (Erichsen, 2016).

A challenge during the New York building import was the underestimated complexity of the import job. It started as a community import where every user in OSM could contribute. Due to the underestimated complexity of the review and upload tasks, and time spent training and supporting new individuals, the team loosely formed a group around the import (Barth, 2014a). The company Mapbox participated in the import with experienced team members. The Mapbox members outpaced the local volunteers by a huge factor (Barth, 2014a). Barth (2014a) writes that the tasks was not hard but demanded a certain learning curve which meant time someone had to spend teaching new volunteers. Having few experts doing the micro-tasks was evaluated as more effective than involving the whole OSM community. The Los Angeles building import allowed everyone to work on the micro-tasks. They developed the Tasking Manager 2, adding new features to fit their needs. Over 100 volunteers contributed to the import job, working on the micro-tasks posted on the Tasking Manager 2 platform (Sambale, 2016). February 2017, all three million buildings were successfully imported into OSM (contributors, 2017). When examining micro-tasks in both import projects, the overall micro-task in Los Angeles contained fewer buildings than in New York's micro-tasks. Erichsen (2016) claim that the micro-tasking method is the best-known method when importing large datasets into OSM. This gives a good indication of how useful and functional this method is.

2.4.5 Task complexity

Task complexity has been identified as one of the most important task properties in a variety of fields studying the relationship between human and computers (Yang et al., 2016). The content of the task reflects its complexity, as well as its attributes (i.e., title and description) and visual features (i.e., the layout and color palette) (Yang et al., 2016). By dividing a larger task into smaller tasks, the complexity is reduces, which is one of the advantages with the micro-tasking method. Yang et al. (2016) claim that complexity reflects the real mental effort that humans need to put into the completion of tasks.

Cognitive load theory refers to the total amount of mental effort being used in the

2. BACKGROUND

working memory. Working memory is determined by the number of information elements that need to be processed simultaneously within a certain amount of time (Barrouillet et al., 2007). A heavy cognitive load can have disadvantageous effects on task completion. Being able to measure and predict task complexity can be highly beneficial for both micro-tasking workers and the requesters (Yang et al., 2016). Too complex tasks can reduce the quality of the results. It is stated that the working memory has a limited capacity of seven plus or minus two elements (or chunks) of information when merely holding information and even fewer (ca. four) when processing information (Leppink et al., 2014).

The cognitive load that is imposed by a task is much higher for beginners than for more advanced students (Leppink et al., 2014). Lack of prior knowledge of how to solve that type of problem forces humans to resort to weak problem-solving strategies (Leppink et al., 2014). Gadiraju et al. (2015b) looked at how training sections affected the micro-tasking worker's performance. The results showed an improved task performance up to 5% and task completion time up to 41% faster than with no training.

2.4.6 Micro-tasking limitations

Getting enough people to use a micro-tasking platform is crucial for its success. Most of the platforms mentioned in this chapter give payments to the workers. Another option is to make the platform as a game, which is also shown in this chapter. Creating a micro-tasking platform without payments or gamification factors the page is likely to have a short life. An exception to this rule is the tasking manager, supported by HOT.

A problem when combining machines and humans is that machines can do their operations in real-time, while humans are unpredictable, they can come and go as they wish. This creates a gap where the micro-tasking platforms cannot guarantee on the task completion time (Difallah et al., 2016).

The human computation abilities can also be overestimated. During the classification of swimming pools in Australia, the Tomnod team faced some unexpected challenges. As described in section 2.4, they used the crowd to classify if a polygon contained a swimming pool or not. When reviewing a random sample from the result, they found an indication that 26% of polygons that contained a pool were identified as not containing pools by the crowd (Kostas, 2016). Further studies also showed that the guilty part was the crowd, because the algorithm had correctly detected polygons containing pools. In a case where the algorithm was 85% confident that the polygon contained a pool, only one voted 'yes', six voted 'no', this polygon do not contain a pool. The solution was to combine the human verdict with the machine's prediction. This example shows that it is important to use the right combination of humans and machines. Tasks that at first seems simple for humans may be more challenging than expected. Basic object detection using machine learning perform very well when used together with human operations.

It is important that the tasks added to a micro-tasking platform consider the tal-

ents and limitations of human workers (Franklin et al., 2011). By using knowledge provided in Erichsen (2016), from OpenStreetMap's usage of the micro-tasking, this thesis will examine and hopefully reveal the limitations and talents of human workers when dealing with geospatial micro-tasks.

3 | Experiment and Web application

This chapter will introduce the developed experiment and the implemented web-application hosting the experiment. The web application containing the experiment that captures the data for the analyses. A pilot-test of the web application is used to secure the usability and discover weaknesses in the implementation. How the pilot-test was executed and some preliminary results, is also presented in this chapter. At last, there is a section about how to determine the sample size.

3.1 Experiment

In the experiment, the participants answer two questions, representing two micro-tasks, on three different tasks. The first question asks the participant to click on the footprint layer that best fits the shape of the building shown on the base map. In the second question, it asks the participant to click on the row that gives the most informative information about an arbitrary building. The three tasks in the experiment will ask the same two questions, but each task varies the number of elements the participant has to answer before the micro-task (question) is complete. This section will introduce the two micro-tasks and the three tasks. It will also describe how the building footprints, used in question one, was developed.

3.1.1 The two micro-tasks

The first question (micro-task) asks the participant to click on the color that fits the shape of the marked building(s) on the map best. Here the participant is given two footprint layers covering a building. The participant needs to determine which of the footprints that fit the shape of the building shown on the base map best. This task is highly relevant during building imports and data validation if one has two overlapping geometry layers. This question is inspired by the micro-tasks given in the building imports from section 2.4.4. The import process can not be done fully-automatic in OpenStreetMap, and a human task can be to validate which footprint fits the building shape best. Question one displayed on the web application is presented in figure 3.1. In this example, the participants have three buildings to select before the micro-task is completed.

3. EXPERIMENT AND WEB APPLICATION

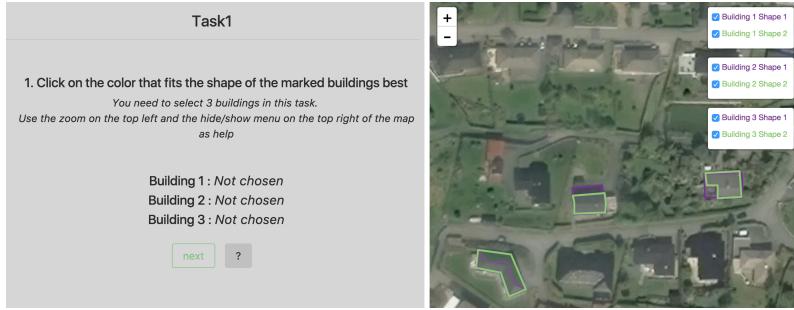


Figure 3.1: Question one as it is displayed in the web application during task B

The second question (micro-task) asks the participant to select the most informative row(s) that describe random buildings best. Question two is an interpretation task where the participant needs to interpret the information written in the table to decide which row(s) gives the most informative information about an arbitrary building. This task is not easy to solve automatic through a script. The correct answer will vary between buildings and what information is present. In figure 3.2 question two is displayed in the web application. In this example, the participant has to select three rows from the table before the micro-task is completed.

Task1			
2. Select the 3 most informative rows that describes random buildings best			
Each row represents a new building. Think that the information should be informative for everyone, independent of education, background etc.			
Choose	Info 1	Info 2	Info 3
<input type="checkbox"/>	Height: 9 m	Gnr: 33	Bnr: 169
<input type="checkbox"/>	Country: Norway	City: Bergen	Address: Hammarslandgrenda 66
<input type="checkbox"/>	Validation date: 20160816	Registered: Yes	Area: 1015,9
<input type="checkbox"/>	Building type: Detached house	Building levels: 3	Building material: Brick and wood
<input type="checkbox"/>	Address: haMarslaNgrenda	Municipality: Bergen	Country: Unknown
<input type="checkbox"/>	Source: Photogrammetric data capture	Building: House	Amenity: Place of residence

Figure 3.2: Question two as it is displayed in the web application during task B

3.1.2 The three experiment tasks

The experiment consists of three different tasks, in addition to a training task. It will explore how the amount of workload demanded in each task influence the task performance. In this thesis, the number of task elements in each micro-task will determine the amount of workload. The task complexity and workload is dependent on mental effort and cognitive load, as mentioned in section 2.4.5. Each task contains six elements, but the tasks vary how many elements that are necessary to answer before the micro-task is completed. One task will serve the participant with micro-tasks

containing one element, called task A. This task demands the smallest cognitive load and the lowest complexity. The next task will serve the participant with micro-tasks containing three elements, called task B. This number is just below the limit of how much information humans can process (Mandler, 2013). The last task will serve the participant with micro-tasks containing all six elements, called task C. This number exceeds the human capacity when processing information according to Leppink et al. (2014). An illustration of the tasks is shown in figure 3.3.

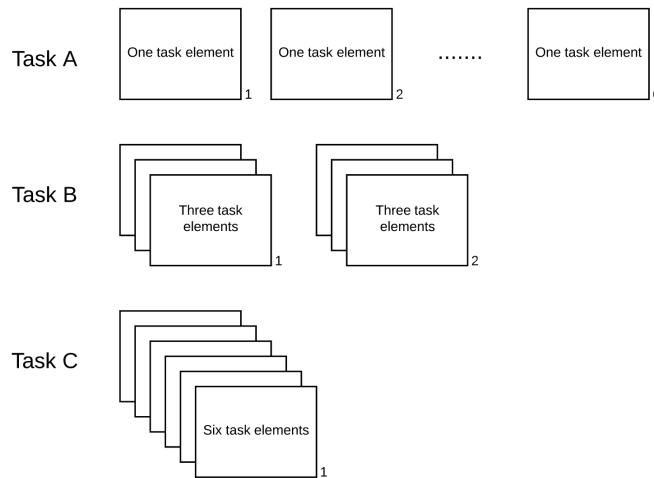


Figure 3.3: Illustration of the three tasks given in the experiment

By choosing three elements in one task and six elements in the other task, this thesis can determine if the theories about the limited capacity of the human brain also apply to micro-tasks involving geospatial data. One task will only contain one element as a minimum cognitive load task. The three tasks can help answer how many elements a human can process at the same time without impacting the quality of the result. The goal is to determine a preferred number of elements to include in a micro-task. This information can be useful when developing micro-tasks to achieve adequate task progress as well as accurate results.

The “magical” number of four has been demonstrated to limit much of human information processing (Mandler, 2013). It is said that polygon comparison demand medium cognitive load (Kiefer et al., 2016), which is what the participants do in the first question in the experiment. Kiefer et al. (2016) argues that high cognitive load may lead to less efficient map reading and spatial orientation, as well as decreased spatial learning. Since polygon comparison demand medium cognitive load, question one should at least not be too demanding on the one- and the three elements micro-tasks.

A concern is that inexperienced participants will have a larger struggle than the expe-

3. EXPERIMENT AND WEB APPLICATION

rienced participants. By dividing the participants into experienced and inexperienced categories, the results from the experiment can help determine if geospatial micro-tasks are too demanding for the inexperienced individuals.

3.1.3 Determining the building footprints used in question one

According to Fan et al. (2014), there was over 77 million buildings in the OpenStreetMap (OSM) database in 2013. Today, over 200 million objects have the key-value building¹, which in OSM is used to mark areas as buildings. A study of the geometries of building footprints in the city Munich reveal a huge diversity in the geometries (Fan et al., 2014). Fan et al. (2014) used four criterion: *completeness*, *semantic accuracy*, *position accuracy* and *shape accuracy*, to evaluate the quality of the building footprints in OSM. In the creation of the two footprints representing the buildings used in question one, the quality criterion shape- and position accuracy was emphasized. The goal is to create shapes that match realistic cases that occur for instance in OSM.

Shape accuracy evaluates how well the layer matches the building in an aerial image. Fan et al. (2014) mentions two main reasons to why building footprints are simplified in OSM. The first reason is the difficulties following building details when looking from a bird's eye view. The second reason is the limited resolution of the Bing aerial images used in OSM during digitalization. In question one, two footprints are drawn with one of them matching the building shape better than the other. The participant has to use a base map to determine which layer fits the building best. This will test if the participants manage to make correct shape judgments by only using a base map as a reference.

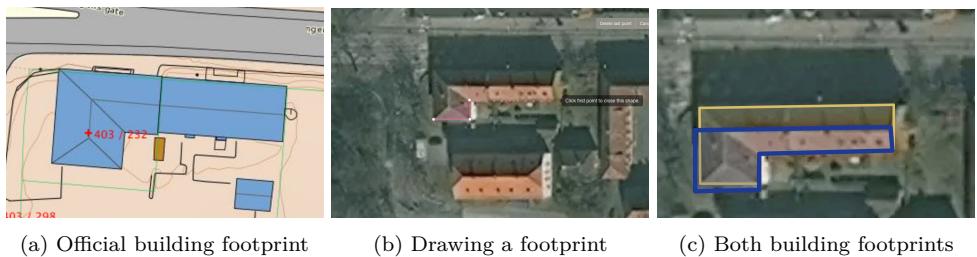


Figure 3.4: Creating of footprint layers used in question one

Position accuracy evaluates how well the coordinate value of a building relates to the reality on the ground. Fan et al. (2014) tested the accuracy of buildings in OSM and concluded with a mean offset of 4.13 m. The low positional accuracy of OSM building footprints data is caused by the limited resolution of Bing map images. By combining shape- and position accuracy in some of the cases used in question one, this study

¹The number of objects with the key-value building was found on the tag info page; <https://taginfo.openstreetmap.org/keys/building>

can also determine if participants manage to evaluate both factors. In this study, the participants do not have available information about what the true ground coordinates are. Therefore position accuracy will be examined by shifting one of the layers. The correct positional accuracy will be at the building in the aerial image.

3.2 Web application

This thesis used an online web-based survey to conduct the experiment. An online survey avoids the cost and effort of printing, distributing, and collecting paper forms. Many people prefer to answer a brief survey displayed on a screen instead of filling in and returning a printed form (Ben and Plaisant, 2009). The participants do not have to share the same geographic location as the researcher.

An online web environment also makes it easier to use an interactive map which is necessary to answer question one. It is not possible to have an experiment involving interactive maps on a piece of paper. A web application is also the obvious way of implementing maps that can handle user interaction. Making it online, available via URL, makes the distribution faster and easier. This also makes the experiment more realistic since the distribution of micro-tasks is through platforms accessible on the web, shown in section 2.4.1.

A common web programming language is JavaScript, together with the library React² creates the client side of this thesis application. The client communicates with a server that fetches the task elements and saves the task results in a PostGIS database. The server is written in Python with the framework Django³. The PostGIS database contains the task elements and the task results gathered from the participants.

3.2.1 React application

The React application was created to serve the experiment to all participants. It contains all the steps of the experiment. First, the participant registers, giving information about age, gender and answers yes or no on the following two questions: 1) "Do you have experience of working with geospatial data?", 2) "Have you heard of micro-tasking before?". Next, the participant is given an introduction page with a detailed introduction video describing how to answer the two questions, and how to interact with the map and building layers. A training task comes after the instruction video. In the training task, the participant solves both questions just like the normal tasks, but it contains different building footprints and two elements to not replicate the other tasks. After the training task, the participant starts with the experiment containing the three tasks, with a short survey after each task. The survey asks the participant to rate the task difficulty between one and five, and if the participant tried

²React is an open-source JavaScript library for building user interfaces. In React, the displayed data can change without reloading the page. Its main goal is to be fast, simple and scalable.

³Django is an open-source web framework written in Python. Its primary goal is to ease the creation of complex, database-driven websites.

3. EXPERIMENT AND WEB APPLICATION

his/hers best or was interrupted during the task. The participant can also write a comment. The interfaces for the two questions used in each task is shown in figure 3.1 and 3.2 in section 3.1.1. The registration form and survey interface is shown in figure 3.5.

(a) Registration site

(b) After each task survey site

Figure 3.5: Interface for the client application

To ensure random and independent observations multiple measurements were implemented during the development of the React application. The main measurements to ensure random, independent observations was:

- Random order on the three tasks
- Random building footprint pairs in the tasks
- Random color on the building footprints
- The two building footprints was drawn on the map in a random order
- Random order on where the information was positioned in the table

When distributing an online experiment, the result can be inconsistent since the researcher is not present to either control the participant or the environment surrounding the participant. The researcher implemented functions securing the completeness of the data. The buttons navigating to the next question and task was disabled until enough building footprints or rows were selected. The participant could not submit their task result before answering both questions completely. The submit button on the register form and after each task survey only submitted the answers if all fields were answered. If a field was missing, a message appeared, asking the participant to fill out the entire form before submitting. All saved task results were complete in the database thanks to these measurements.

3.2.2 Data acquisition

In this study, the independent variables are: 1) experienced or inexperienced participant, 2) number of elements in the task, 3) age and 4) gender of the participant. Independent variables are factors we think might influence the results of the study

(Kitchin and Tate, 2000, p. 49). When a participant registers at the start of the survey, the independent variables are generated. Participants who answer yes to the question "Do you have experience of working with geospatial data?" are registered as experienced. The independent variables are believed to influence the dependent variables. Dependent variables are factors the study is interested in explaining (Kitchin and Tate, 2000, p. 49). In this study, the dependent variables are: 1) time spent on each question and task, 2) the number of correctly chosen elements in each question and task, and 3) how difficult the participant thought the task was. The variables are listed in table 3.1.

Independent variables
1) Experienced or inexperienced participants
2) Number of elements in the task
3) Age
4) Gender
Dependent variables
1) Time spent on each question and task
2) Number of correctly chosen elements in each question and task
3) Task difficulty

Table 3.1: The experiments independent and dependent variables

The React application generates and saves the task results. In both questions, the time and number of correct elements are registered. The React application has a timer that registers time elapsed on both questions and adds the time measurements together to get the total time spent on the task. The time only reflects how long the participants spent solving the two questions. Time spent loading new layers, moving to the next question, etc., is not included. The building footprints and rows the participant selected is also registered. Before saving the result it counts the number of correctly chosen elements in each question and then adds the number together to get the total number of correctly chosen elements in the task. A participant can maximum have twelve correct elements, six from each question. Total time and number of correct elements are the two primary dependent variables, and they create the basis of the statistical analyses together with the independent variables mentioned above. Participants task results are saved after each task and contain the following information:

- Task number (which task)
- Task order number (which order)
- Time spent on question one
- Time spent on question two
- Total time spent on the task
- Correct elements in question one

3. EXPERIMENT AND WEB APPLICATION

- Correct elements in question two
- Total correct elements in the task

3.3 Pilot test

Testing the experiment before actual use is highly recommended (Ben and Plaisant, 2009). A pilot test provides an opportunity to validate the wording of the tasks. It also helps understand the time necessary for completing the survey, which should be communicated to the participants (Schade, 2015). The pilot test was carried out on a small sample of users. Results from the pilot test in this thesis was used to make improvements to the actual survey, to the react application and to find errors or weaknesses in the database models.

After the pilot test, the usability was measured. Usability in this thesis was measured with the *System Usability Scale*(SUS) because it gives a subjective measure of usability. The *System Usability Scale* questionnaire consists of ten statements where the participants rate their agreement on a five-point scale (Ben and Plaisant, 2009). SUS was developed to be quick and straightforward, but also reliable enough to be able to compare performance changes between versions (Brooke, 1996). It is also easy to administer the participants through the usability test, and it can be used on small sample sizes and still give reliable results (Affairs, 2013).

The usability is important to measure. If the participants do not understand how the web application works, they will probably not do the survey since they have to invest time in understanding what to do. It is also important to get enough participants to do the whole survey and not quit halfway in frustration of not understanding it properly. The *System Usability Scale* can effectively differentiate between usable and unusable systems (Affairs, 2013).

3.3.1 Execution of the pilot test

The pilot test was conducted with a total of eight participants, five experienced and three inexperienced participants aged from 22 to 64 years. The test started with a brief information about this study and the experiment. They were told to "talk out load" during the test and no help or guidance was given to the participants. The participants was observed while they conducted the survey. After the survey a *System Usability Scale* questionnaire was answered by the participants. In the end, the participants were asked to give general feedback on the web application. The SUS score and the feedback were then used to determine the usability of the React application and to determine which improvements to be done.

3.3.2 Results from the pilot test

The average SUS score was 84.64 out of 100. Anything above 68 is considered above average (Affairs, 2013). When adding the SUS score to an adjective rating, a score of 85.5 or higher is described as excellent (Bangor et al., 2009). A score of 84.64 is then described as good/excellent. This result gives a strong indication that the React application is user-friendly.

All participants thought that the instruction video was confusing. It was short, the instructions went too fast, and it missed voice descriptions. The instructions needed major improvements, which was an important discovery. The purpose of the video is to give the participant an introduction to how to answer the two micro-tasks. It should include important instructions, particularly useful for participants not used to working with interactive maps.

The participants struggled to use the layer control because of missing layer labels on the map. In the layer control, the buildings are numbered, so it was not possible to see which layer footprint was, i.e., building one. They also struggled with determining when a building footprint was selected or not. The lack of labels on the map layer was done on purpose to get the task as realistic as possible. The process of selecting the best fitting building footprint needed improvements. It had to be clearer that one had to click on the layer on the map to choose a footprint, not by using the layer control as some thought. This part was added to the video with voice description, describing in detail how a footprint was selected.

The test data was used to find errors or weaknesses in the database model. The data was extracted from the PostGIS database and saved in CSV files. There were a few errors and weaknesses found during the statistical tests. Changes to the database models were necessary. The task result model was improved by adding four new fields. The additional fields will mainly help with creating plots to interpret the data better and to visualize the different results more easily.

The average time spent on the survey was 18 minutes. The two slowest participants used on average 33 minutes, while the rest of the participants spent on average 13 minutes to complete the survey.

In the pilot test, the same building footprints (question one) and information rows (question two) were used in all three tasks. At the end of the pilot test, the author asked the participants if they remembered the buildings and meta information from the previous tasks. Seven out of eight answered yes on the question. This information was valuable. If every participant does a better job at the last task, because they remember the elements from earlier tasks, the result will not be as useful. This almost matches the number of participants who remembered the previous elements in the last task. This finding made the author create three different task element groups. Each task will then contain new building footprints and meta information. The three element groups were randomly assigned to each task, to avoid the three building groups influencing the result. The risk of participants remembering previous task elements disappeared with this decision.

3.4 Determining sample size

The sample size is influenced by various factors, including the purpose of the study, population size, the risk of selecting a "bad" sample and the allowable sampling error (Israel, 1992).

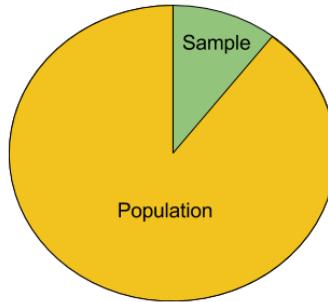


Figure 3.6: Population vs. sample

A sample is a collection of observations and is the subset of a population, illustrated in figure 3.6. The population size in this survey is not easily determined. A population is the collection of individuals of a particular type (Walpole et al., 2012). All individuals with access to a computer and the internet interested in contributing to micro-tasks can be one description of the population. It is important that the sampled population and the target population is similar to one another.

There are three possible ways of determining the sample size in this study. The first option is to use a sample size from a similar study. The risk is to repeat errors that were made in determining the sample for another study. The second option is to rely on published tables, depending on precision, confidence levels, and variability. According to table 1 in the Israel (1992) paper, an accuracy of 0.05, confidence interval of 95% and a population size greater than 100'000, the necessary sample size is 400. If the accuracy is changed to 0.1, the sample size necessary decreases to 100 (Israel, 1992). The numbers found in the table must reflect the number of obtained responses. The last approach is to use formulas to calculate the sample size. The formulas require the standard deviation and how much variance to expect in the response (Israel, 1992; Smith, 2013). Israel (1992) mentions that table 1 (in his paper) gives a useful guide for determining the sample size and that formulas are used if the study has a different combination of precision and confidence. This study will use the table result since the combinations match this study.

It is important to mention that the quality of the sample is as important as the size. The more variable the sampled data is, the larger the sample size is required (Israel, 1992). It is also desirable to choose a random sample, which means that the observations are made independently and arbitrary. The main purpose of using a random sample is to obtain correct information about the unknown population parameters (Walpole et al., 2012).

4 | Results

This chapter will present the result of the conducted experiment. It will summarize who the participants were, the gathered data and then present the statistical results from the analysis performed on the gathered data. It took eight days to collect enough participants completing the experiment. Only 38% of the registered participants completed the experiment. All task results saved in the database was valid and could be used in the analyses. The analyses were calculated in Python, using statistical packages as SciPy, Numpy, and Pandas. The data was extracted from the database using Django Queryset and saved in CSV files. The implemented statistical methods are available on GitHub ¹.

4.1 Participants

One of the benefits of using an online based experiment is the potential to reach a huge number of people. The challenge is how to reach out to people and make them aware of the existence of the experiment. The solution was to use mailing lists and sites available to the author. All students at *Engineering and ICT* was emailed, as well as a mailing list reaching out to the Norwegian OpenStreetMap community. The web application was also published at the Geoforum website (www.geoforum.no) and Facebook page. Geoforum is a Norwegian association for individuals and companies working in the field of geomatics. The participants receiving the email or looking at the Geoforum site could click on the web application URL and access the experiment from there.

After eight days there was in total 461 task results in the database. 402 participants registered on the website during the data gathering period, but only 38% of the registered participants completed all three tasks. This number was surprisingly low, but time to complete the whole experiment can probably explain why so few completed. It probably lasted too long for participants to have the patience to finish. 152 participants completed all three tasks. Results from participants not completing all three tasks are also included in the dataset. Including these result is not a problem. The three tasks were given to the participant in a random order, and the tasks are also independent, containing different buildings and rows.

The mean participant age was 31.5 years and the median 25 years. The youngest participant was 19 and the oldest 58 years. 33% of the participants were female and 66% male. The average male was 33.5 years old and the average female 31.2 years old. 19% of the participants that completed the survey said they had heard of micro-tasking before. 53% of the participants stated that they had experience of working with geospatial data. The distribution between experienced and inexperienced participants was approximately even.

¹<https://github.com/annesofie/thesis-statisticmethods>

4. RESULTS

Random and independent observations are essential, and a random task order was used to ensure this. Analyzing the task results and the order the tasks were presented in gives an acceptable result. The distribution of how many times the three tasks (Task A, B and C) occurred as the first, second and last task in the analyzed data is approximately evenly distributed. This is shown in figure 4.1.

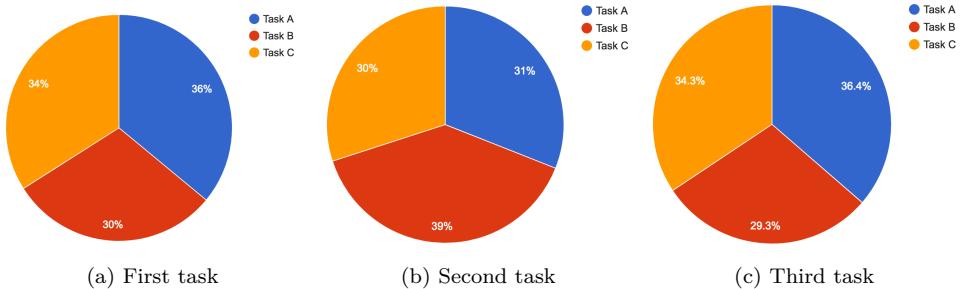


Figure 4.1: The distribution of the task order in the analysed data

4.2 Statistics theory

This section will give an introduction to the statistics used in this thesis. The thesis will examine the data with both parametric and non-parametric methods. A non-parametric method is much more efficient than the parametric procedure when the set of data used in the test deviates significantly from the normal distribution (Walpole et al., 2012). There are also some disadvantages using nonparametric methods. The methods will be less efficient, and to achieve the same power as the corresponding parametric method a larger sample size is required. If parametric and nonparametric tests are both valid on the same set of data, the parametric test should be used (Walpole et al., 2012).

4.2.1 Normal testing

The sampling distribution depend on the distribution of the population, the size of the samples, and the method of choosing the samples (Walpole et al., 2012). Sampling distribution describes the variability of sample averages around the population mean μ . All parametric statistics assumes normally distributed, independent observations. Parametric tests are preferred in statistics because it has more statistical power than non parametric tests (Frost, 2015). The power of a test is the probability of correctly rejecting a false null hypothesis, which in this case is the ability to detect if the sample comes from a non-normal distribution. To determine if a sample is normally distributed there exists both visual methods and normality tests to assess the samples normality. A visual inspection of the sample's distribution is usually unreliable and does not guarantee that the distribution is normal (Pearson et al., 2006). Presenting the data visually gives the reader an opportunity to judge the distribution themselves.

In this thesis histograms are used to visually analyze the data for normality.

Normality tests compare the scores in the sample to a normally distributed set of scores with the same mean and standard deviation (Ghasemi and Zahediasl, 2012). There are multiple normality tests, and deciding which test to use is a complex task. This study needs a test that does not require every value to be unique, in addition to handling ties.

The D'Agostino-Pearson omnibus test stand out as the best choice. This test first computes the skewness, figure 4.2, and kurtosis, figure 4.3, to quantify how far from the normal distribution the sample is from the terms of asymmetry and shape. It calculates how far each of these values differ from the value expected with a normal distribution (Pearson et al., 2006). It works well even if all values are not unique (Motulsky, 2013). The test also works well on both short- and long-tailed distributions (Yap and Sim, 2011).

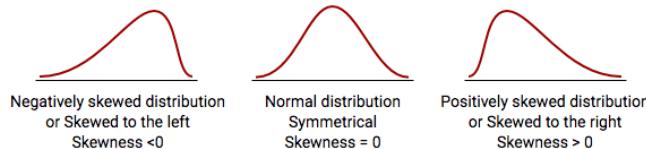


Figure 4.2: Skew (MedCalc Software bvba, 2017)

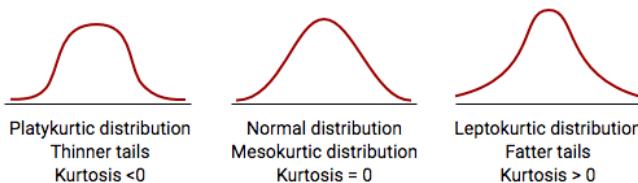


Figure 4.3: Kurtois (MedCalc Software bvba, 2017)

The D'Agostino-Pearson test uses the following hypothesis:

$$H_0: \text{The data follows the normal distribution}$$
$$H_A: \text{The data do not follow the normal distribution}$$

For small sample sizes, normality tests have little power to reject the null hypothesis, and therefore small sample sizes most often pass normality tests. For large sample sizes, significant results would be derived even in the case of a small deviation from normality (Pearson et al., 2006). When the null hypothesis cannot be rejected, there are two possible cases. First case is to accept the null hypothesis and the second case is that the sample size is not large enough to either accept or reject the null hypothesis (The Pennsylvania State University, 2017). An acceptance of the null hypothesis

4. RESULTS

implies that the evidence was insufficient, and the result does not necessarily accept H_0 , but fails to reject H_0 (Walpole et al., 2012). Both visual analysis and D'Agostino-Pearson test will be used to assess the normality assumption in this thesis.

4.2.2 Hypothesis testing

The null- and alternative hypothesis are statements regarding a difference or an effect that occur in the population of the study. The alternative hypothesis (H_A) usually represents the question to be answered or the theory to be tested, while the null hypothesis (H_0) nullifies or opposes H_A (Walpole et al., 2012). The sample collected in the study is used to examine which statement is most likely. When the hypothesis is identified, both null and alternative, the next step is to find evidence and develop a strategy for or against the null hypothesis (Lund Research Ltd, 2013a).

The next step is to determine the level of statistical significance, often expressed as the *p-value*. A statistical test will result in the probability (*the p-value*) of observing your sample results given that the null hypothesis is true. According (Walpole et al., 2012), a significance level widely used in academic research is 0.05 or 0.01. 0.05 significance level will be applied in this thesis analyses.

4.2.2.1 Two-sample t-test

When estimating the difference between two means a two-sample t-test is used (Walpole et al., 2012). A two-sampled test assumes two independent, random samples from distributions with means $[\mu_1, \mu_2]$ and variances $[\sigma_1^2, \sigma_2^2]$. The hypothesis tested on two means can be written as

$$\begin{aligned} H_0: & \text{ Sample means are equal} \\ H_A: & \text{ Sample means differ} \end{aligned}$$

The two-sample t-test is used to estimate if differences between two means are significant. In a two-sample, two-sided, t-test ($\mu_1 - \mu_2 \neq 0$) the null hypothesis is rejected when (Walpole et al., 2012, p. 345):

$$|T| > t_{\frac{\alpha}{2}, v} \quad (4.1)$$

In a two-sample, one-sided, t-test the null hypothesis is rejected when [(Walpole et al., 2012), p. 350]:

$$T > t_{\frac{\alpha}{2}, v} \quad (4.2)$$

$$T < -t_{\frac{\alpha}{2}, v} \quad (4.3)$$

Equation 4.2 is used on one sample test where the alternative test is to check if the

mean is greater than zero ($\mu_1 - \mu_2 > 0$), and the 4.3 equation is used on hypothesis where the test is to check if the mean is lower than zero ($\mu_1 - \mu_2 < 0$). T is the calculated statistical value and t is the critical value with the given significance level (α) and degree of freedom (v). The critical value is found in the table of Critical values for t-distribution.

Before doing tests on the two means, the Levene's Test is used to examine if the samples are from populations with equal variances. It tests the hypothesis:

$$H_0: \text{Input samples are from populations with equal variances}$$
$$H_A: \text{Input samples are from populations that do not have equal variances}$$

If we can assume equal variances in the two samples and the samples are normally distributed, a two-sampled t-test may be used.

The hypotheses in this study, used to answer the research questions, that is tested with a two-sampled t-test (if the conditions mentioned above are valid) is listed in figure 4.4.

1

$$H_0: \text{There is no difference in time spent on the tasks between the participants}$$
$$H_A: \text{Time spent differs between experienced and inexperienced participants}$$

2

$$H_0: \text{Experienced participants do not finish the tasks faster than inexperienced}$$
$$H_A: \text{Experienced participants finish the tasks faster}$$

3

$$H_0: \text{Equal number of correct elements between experienced and inexperienced}$$
$$H_A: \text{Number of correct elements differs between experienced and inexperienced participants}$$

4

$$H_0: \text{Experienced participants do not have more correct elements}$$
$$H_A: \text{Experienced participants have a higher number of correct elements}$$

Figure 4.4: Two-sample t-test hypotheses

4.2.2.2 Analysis-of-Variance

Analysis-of-Variance (ANOVA) is, according to Walpole et al. (2012), a very common procedure used for testing population means. Where a two-sample t-test is restricted to consider no more than two population parameters, ANOVA can test multiple population parameters. A part of the goal of ANOVA is to determine if the differences

4. RESULTS

among the means of two or more samples are what we would expect due to random variation alone, or due to variation beyond merely random effects (Walpole et al., 2012). *ANOVA* assumes normally distributed, independent samples with equal variance. The equal variance assumption will be tested with Levene's Test mentioned in section 4.2.2.1.

One-way *ANOVA* tests the null hypothesis that two or more groups have the same population mean given that the mean is measured on the same factor or variable in all groups (Lund Research Ltd, 2013c). The hypothesis test can be written as:

$$H_0: \mu_1 = \mu_2 = \dots = \mu_k$$
$$H_A: \text{At least two of the means are different}$$

μ equals the group mean and k represents the number of groups. It is important to check that each group are normally distributed (Lund Research Ltd, 2013c). The weakness of one-way *ANOVA* is that it cannot tell which specific groups were significantly different from each other if H_0 is rejected. To be able to determine which group differ, a *post hoc test* is used. The null hypothesis is accepted if:

$$F < f_{\alpha, v_1, v_2} \quad (4.4)$$

F is the calculated statistical value, and f is the critical value with the given degrees of freedom (v_1, v_2) and significance level (α). If the alternative hypothesis is accepted a *post hoc test* is used to make paired comparisons. The test determines which group differs. This thesis will use Tukey's test to determine which group means are significantly different (Walpole et al., 2012, p. 526). Hypotheses in this study that is used to answer the research questions tested in a one-way *ANOVA* analysis (if the conditions mentioned above are valid) is listed in figure 4.5.

1
H_0 : Task time do not differ between the three tasks
H_A : Task time differ between at least two of the tasks
$Variable = time, group = tasks$
2
H_0 : Number of correct elements in each of the three tasks do not differ
H_A : Number of correct elements between at least two of the tasks differs
$Variable = Number \text{ of } correct \text{ elements}, group = tasks$

Figure 4.5: One-way ANOVA hypotheses

4.2.2.3 Mann-Whitney U test

The Mann-Whitney U test is used to compare differences between two independent groups. This test can be used to conclude whether two populations differ. It can for

instance test if there are differences in medians between groups (Lund Research Ltd, 2013b). In contrast to the t-test, it compares the median scores of two samples instead of the mean score. The test is non-parametric and can therefore be used on samples that are not normally distributed. The test assumes that the samples come from populations with equal variances. If there are ties (identical observations) in the sample a Mann-Whitney U test is preferred (The Scipy community, 2017). When comparing two sample medians the two independent variables (i.e experienced and inexperienced participants) has to have a similar shape. It can test the hypothesis:

$$H_0: \text{The two populations are equal}$$

$$H_A: \text{The two populations are not equal}$$

The null hypothesis is rejected if (LaMorte, 2017):

$$U \leq CriticalValue \quad (4.5)$$

The critical value is found in the table of Critical Values for U and depends on the sample sizes, n_1 and n_2 , and the significant level α . U is the statistical value calculated.

4.2.2.4 Kruskal-Wallis test

The Kruskal-Wallis test is a nonparametric alternative to one-way *ANOVA* (see subsection 4.2.2.2) (Walpole et al., 2012). This test should be used if the assumption of normal distribution failed. As mentioned in this sections introduction, a non parametric method does not assume normality. This test is a generalization of the rank-sum test when there are more than two samples.

Kruskal-Wallis is used to test equality of means in one-way *ANOVA*, so the hypothesis for the Kruskal-Wallis test is:

$$H_0: \mu_1 = \mu_2 = \dots = \mu_k$$

$$H_A: \text{Minimum two of the } \mu_k \text{'s are different}$$

Here μ_k is the rank mean for the group k. The number of observations in the smallest sample is assigned to n_1 , the second smallest to n_2 and the largest sample is assigned to n_k .

The null hypothesis is accepted if (Walpole et al., 2012, p. 668):

$$H < \chi^2_{\alpha} \quad (4.6)$$

4.3 Survey results

This section will first introduce the gathered data and divide the data into samples. Key-values for each sample is listed in tables. The section will test each sample if they follow a normal distribution using D'Agostino and Pearson test and equal variance assumption using Levene's test. When the assumptions in each sample are tested, they are used to answer hypothesis leading to an explanation for the research questions listed in the introduction. After each subsection, a table containing a summary of the test results is presented. This will hopefully make it easier to get an overview of which analyses is completed in the subsection.

4.3.1 Gathered data

The gathered data is analyzed on the two dependent variables: 1) total time used to complete each task and 2) number of correctly chosen elements per task. The gathered data is also divided by the two independent variable pairs: 1) experienced- and inexperienced participants and 2) the three experiment tasks (section 3.1.2). Different combinations of the dependent and independent variables create the foundation of the 22 different samples introduced in this chapter. The sample means \bar{x} , sample median, the standard deviation of \bar{x} , standard error of \bar{x} , minimum and maximum in each sample is included. The sample ID is also listed in the tables and is referred to in the analysis of the data, to easier distinguish which sample is used in which analysis. In all the samples, results from the training task and from participants that were disturbed during the task are removed. We examine four different divisions of the gathered data in the following section.

4.3.1.1 All experienced and inexperienced participants

Table 4.1 and 4.2 are samples containing task results from all experienced and inexperienced participants. The result is grouped by the two dependent variables, total time and the number of correctly chosen elements.

Sample ID	All	Experienced 1	Inexperienced 2
Number of observations	429	229	200
Sample mean \bar{x}	170.32	177.65	161.94
Sample median	155.00	158.00	154.00
Standard deviation of \bar{x}	82.19	88.24	73.99
Standard error of \bar{x}	3.98	5.83	5.23
Minimum in sample	38.00	52.00	38.00
Maximum in sample	657.00	657.00	529.00

Table 4.1: Total time spent on each task (seconds)

Sample ID	All 3	Experienced 3	Inexperienced 4
Number of observations	429	229	200
Sample mean \bar{x}	9.82	9.81	9.83
Sample median	10.00	10.00	10.00
Standard deviation of \bar{x}	1.52	1.53	1.51
Standard error of \bar{x}	0.07	0.10	0.11
Minimum in sample	4.00	5.00	4.00
Maximum in sample	12.00	12.00	12.00

Table 4.2: Number of correctly chosen elements per task

4.3.1.2 All participants, divided by task

In table 4.3 and 4.4 the task results are divided by the three different tasks, grouped by the dependent variables. Task A is the task that used micro-tasks containing one element. Task B is the task with micro-tasks containing three elements, and task C gave all six elements at the same time.

Sample ID	Task A 5	Task B 6	Task C 7
Number of observations	146	142	141
Sample mean \bar{x}	166.38	172.25	172.48
Sample median	150.00	155.50	157.00
Standard deviation of \bar{x}	84.57	84.21	77.95
Standard error of \bar{x}	7.00	7.07	6.56
Minimum in sample	47	50	38
Maximum in sample	657	492	529

Table 4.3: Total time divided by task (seconds)

Sample ID	Task A 8	Task B 9	Task C 10
Number of observations	146	142	141
Sample mean \bar{x}	10.19	9.71	9.55
Sample median	11.00	10.00	10.00
Standard deviation of \bar{x}	1.43	1.53	1.52
Standard error of \bar{x}	0.12	0.13	0.13
Minimum in sample	5.00	5.00	4.00
Maximum in sample	12.00	12.00	12.00

Table 4.4: Number of correctly chosen elements divided by task

4. RESULTS

4.3.1.3 Experienced participants, divided by task

In the tables in this section, only task results from experienced participants are included, and the result is also divided by the three tasks, grouped by the dependent variables.

Sample ID	Task A 11	Task B 12	Task C 13
Number of observations	77	80	77
Sample mean \bar{x}	173.04	176.70	181.06
Sample median	158.00	156.00	165.00
Standard deviation of \bar{x}	96.76	86.13	79.70
Standard error of \bar{x}	11.03	9.63	9.08
Minimum in sample	57.00	52.00	53.00
Maximum in sample	657.00	492.00	463.00

Table 4.5: Total time per task for experienced participants, divided by task (seconds)

Sample ID	Task A 14	Task B 15	Task C 16
Number of observations	77	80	77
Sample mean \bar{x}	10.29	9.66	9.48
Sample median	11.00	10.00	10.00
Standard deviation of \bar{x}	1.32	1.64	1.47
Standard error of \bar{x}	0.15	0.18	0.17
Minimum in sample	7.00	5.00	5.00
Maximum in sample	12.00	12.00	12.00

Table 4.6: Number of correct elements per task for experienced participants, divided by task

4.3.1.4 Inexperienced participants, divided by task

In this section, the task results from only inexperienced participants are included, and the result is also divided into the three survey tasks. Table 4.7 is the total time variable and 4.8 the number of correctly chosen elements.

Sample ID	Task A	Task B	Task C
	17	18	19
Number of observations	71	64	65
Sample mean \bar{x}	158.30	165.69	162.23
Sample median	148.00	154.50	154.00
Standard deviation of \bar{x}	67.57	80.93	74.53
Standard error of \bar{x}	8.02	10.12	9.24
Minimum in sample	47.00	50.00	38.00
Maximum in sample	487.00	455.00	529.00

Table 4.7: Total time per task for inexperienced participants, divided by task (seconds)

Sample ID	Task A	Task B	Task C
	20	21	22
Number of observations	71	64	65
Sample mean \bar{x}	10.07	9.78	9.61
Sample median	10.00	10.00	10.00
Standard deviation of \bar{x}	1.54	1.38	1.57
Standard error of \bar{x}	0.18	0.17	0.19
Minimum in sample	5.00	6.00	4.00
Maximum in sample	12.00	12.00	12.00

Table 4.8: Number of correct elements per task for inexperienced participants, divided by task

4.3.2 Normality tests

The samples need to be tested to see if they follow a normal distribution. This test is important. It determines if a parametric or a non-parametric test should be used when including the different samples in the hypothesis tests. The section about normal testing (4.2.1) concluded that the D'Agostino and Pearson normality test should be used in this thesis. A visual interpretation of histograms will also be a part of the normality test. D'Agostino-Pearson uses the following hypothesis:

$$H_0: \text{The data follows the normal distribution}$$

$$H_A: \text{The data does not follow the normal distribution}$$

4.3.2.1 All participants, total time variable

This section will test if sample 1 and 2 (table 4.1) follow a normal distribution. Sample 1 and 2 will be used to determine if there are a significant difference in time spent on the tasks between experienced and inexperienced participants.

4. RESULTS

A visual interpretation of histogram 4.6a and 4.6b gives an indication that sample 1 and 2 does not follow the normal distribution. Both histograms are positively skewed (figure 4.2). Samples involving time measurements are rarely normally distributed. This is because the samples will always be skewed since it is impossible to have negative time and there will always be a limit to how fast a participant can finish the task.

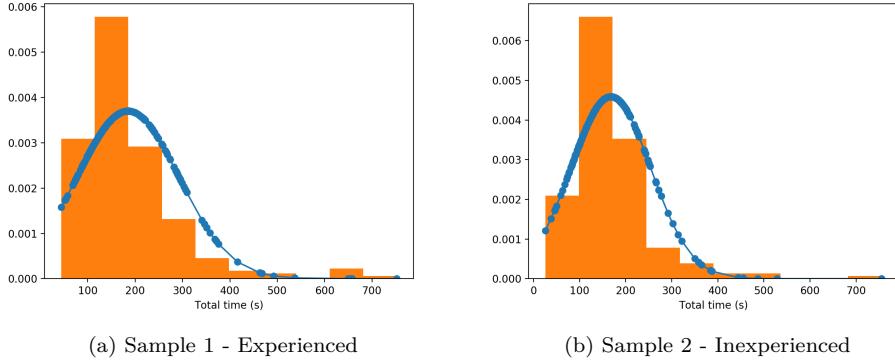


Figure 4.6: Histograms with normal distribution fit with samples containing total time to complete each task

D'Agostino and Pearson tests confirmed the visual interpretation with a significance level of 5%. Both samples obtained p-values lower than the significance level, and the null hypothesis is rejected. Sample 1 and 2 do not follow the normal distribution with a confidence level of 95%.

D'Agostino and Pearson normality test
Significance level: 5%

Sample 1

P-value: $3.874 * 10^{-22}$

The p-value is lower than the significance level (0.05), the null hypothesis is rejected and H_A accepted.

Sample 2

P-value: $2.574 * 10^{-21}$

The p-value is lower than the significance level (0.05), the null hypothesis is rejected and H_A accepted.

In both sample 1 and 2, the p-value was significantly lower than the significance level of 5%. Data transformations are commonly used tools to improve the normality of

a samples distributions, but there are many types of data transformations. Osborne (2010) claim that almost all tests, even non-parametric tests, benefit from improving the normality of the samples, especially when the normality test is significantly denied. Typical traditional transformations are square root, inverse or converting to logarithmic scales (Osborne, 2010).

A Box-Cox power transformation (Box-Cox) is used in this thesis. This transformation can only be used on positive data. The data gathered in this thesis will never be below zero, so this is not a concern. Box-Cox takes the idea of having a range of power transformations (i.e., square root $x^{\frac{1}{2}}$, inverse x^{-1}) available to improve the effectiveness of normalizing and variance equalizing for both positively- and negatively-skewed variables (Osborne, 2010). This transformation will always use the appropriate conversion to be maximum effective in moving each sampled data towards normality. This is the reason why this thesis will use the Box-Cox transformation.

Sample 1 and 2 after a Box-Cox power transformation is shown in histogram 4.7a and 4.7b. A visual inspection gives a good indication that the transformed data follows a normal distribution after the transformation. The skewness looks approximately zero.

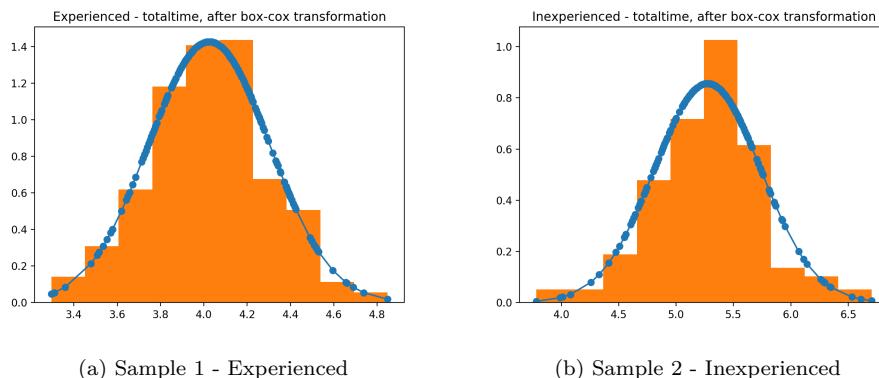


Figure 4.7: Histograms with normal distribution fit after Box-Cox power transformation

The transformed data is applied to the D'Agostino and Pearson test. This test confirms the visual analysis, since both sample 1 and sample 2 follow a normal distribution after the Box-Cox with a confidence level of 95%. Sample 2 passes the test, even though the p-value (0.06) is very close to the significance level (0.05). The calculated p-value is larger than the significance level of 5%.

4. RESULTS

D'Agostino and Pearson normality test

(After Box-Cox transformation)

Significance level: 5%

Sample 1: Experienced, total time per task

P-value: 0.849

The p-value is higher than the significance level (0.05), the null hypothesis is accepted.

Sample 2: Inexperienced, total time per task

P-value: 0.0623

The p-value is higher than the significance level (0.05), the null hypothesis is accepted.

The assumption that sample 1 and sample 2 follows a normal distribution is now accepted and the transformed data can be used in parametric methods.

4.3.2.2 Inexperienced participants divided by task, total time variable

This section will test if sample 17, 18 and 19, table 4.7, follows the normal distribution. These samples will be used to test whether there is a significant difference in total time between the three tasks when only looking at inexperienced participants.

A visual analysis of the histograms A.9a, A.9b and A.9c show a positive skewness. The skew is less than the histograms in figure A.6, but is most likely too large for the samples to be normally distributed.

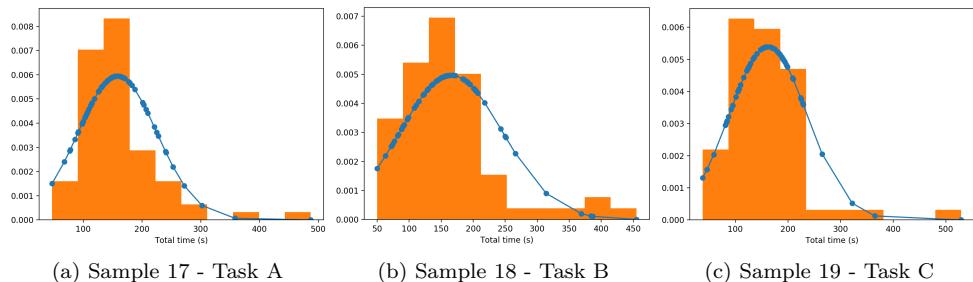


Figure 4.8: Histogram with normal distribution fit

The D'Agostino and Pearson normality test agrees with the visual analysis. The three obtained p-values (1.586×10^{-11} , 1.773×10^{-6} and 2.312×10^{-11}) are all significantly lower than the significance level of 0.05. Sample 17, 18, and 19 do not follow a normal distribution with a confidence interval of 95%. The null hypothesis is rejected.

The samples are Box-Cox power transformed. The histograms after transformation (A.10a, A.10b, and A.10c) are visually evaluated to be normally distributed.

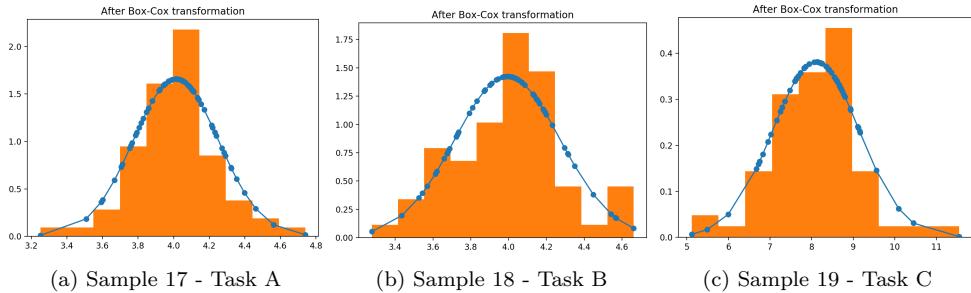


Figure 4.9: Histogram with normal distribution fit after Box-Cox

A new D'Agostino and Pearson normality test on the transformed data confirms that all three samples are drawn from a normal distribution with a significance level of 5%. The obtained p-values (0.139, 0.909 and 0.067) are larger than 0.05, and the null hypothesis is accepted. These samples can be used in parametric methods

The calculations and visual inspections of the rest of the 21 samples are located in the appendix. Table 4.9 show the results of all normal tests performed with D'Agostino and Pearson method. The appendix column refers to where in the appendix the calculations are located.

4. RESULTS

4.3.2.3 Normality test summary

Table 4.9: Summary of normality tests done in section 4.3.2

	Sample ID	Normally distributed	Normally distributed after Box-Cox power transformation	Appendix
<i>Total time</i>				
Experienced	1	No	Yes	
Inexperienced	2	No	Yes	
<i>Correct elements</i>				A.1
Experienced	3	No	No	
Inexperienced	4	No	No	
<i>Total time</i>				A.2
Task A	5	No	Yes	
Task B	6	No	Yes	
Task C	7	No	Yes	
<i>Correct elements</i>				A.3
Task A	8	No	No	
Task B	9	Yes	Yes	
Task C	10	No	Yes	
<i>Total time, experienced participants</i>				A.4
Task A	11	No	Yes	
Task B	12	No	Yes	
Task C	13	No	Yes	
<i>Correct elements, experienced participants</i>				A.5
Task A	14	Yes	<i>not tested</i>	
Task B	15	Yes	<i>not tested</i>	
Task C	16	Yes	<i>not tested</i>	
<i>Total time, inexperienced participants</i>				A.6
Task A	17	No	Yes	
Task B	18	No	Yes	
Task C	19	No	Yes	
<i>Correct elements, inexperienced participants</i>				A.7
Task A	20	No	Yes	
Task B	21	Yes	Yes	
Task C	22	No	Yes	

4.3.3 Levene's test of Equality of Variance

As mentioned in section 4.2.2.1, 4.2.2.2, and 4.2.2.3, the two sample t-test, one-way ANOVA and Mann-Whitney U test assumes that the samples come from populations with equal variances. This assumption will be examined with Levene's test. The hypothesis tested is:

$$H_0: \text{Input samples are from populations with equal variances}$$

$$H_A: \text{Input samples are from populations that do not have equal variances}$$

The null hypothesis is accepted if the obtained p-value is higher than the significance level. Table 4.10 contains the summary of the Levene's test performed on all sample pairs. All sample pairs accepted the null hypothesis except sample 1 and 2, who obtained a p-value lower than the significance level. The test used a significance level of 5% on all the tests.

Table 4.10: Summary of Levene's tests

Participants		Obtained p-value	Samples come from populations with equal variances
All	<i>Total time</i> Sample 1 and 2 <i>Correct elements</i> Sample 3 and 4 <i>Total time, divided by task</i> Sample 5, 6, and 7 <i>Correct elements, divided by task</i> Sample 8, 9, and 10	0.030 0.823 0.636 0.805	No Yes Yes Yes
Experienced	<i>Total time, divided by task</i> Sample 11, 12, and 13 <i>Correct elements, divided by task</i> Sample 14, 15, and 16	0.972 0.724	Yes Yes
Inexperienced	<i>Total time, divided by task</i> Sample 17, 18, and 19 <i>Correct elements, divided by task</i> Sample 20, 21, and 22	0.499 0.626	Yes Yes

4. RESULTS

4.3.4 Hypothesis testing

This section will test the hypothesis listed in figure 4.4 and 4.5 to answer the three research questions written in the introduction. The order of the tests will be the same as the numbering of the samples. Which statistic method, from the theory section (4.2.2), that is used to answer the hypothesis tests is determined by the results of the normality test (4.3.2) and equal variance test (4.3.3).

4.3.4.1 Differences in total time between experienced and inexperienced participants

This section will test if there is any difference in total time spent on the tasks between experienced and inexperienced participants. The test is covered by sample 1 and sample 2 in table 4.1. Sample 1 is experienced and sample 2 inexperienced participants. Both samples was normally distributed after a Box-Cox transformation (4.3.2.1). A two-sample t-test will be used to answer this hypothesis since the normality assumption is valid. The hypothesis tested in this section is number one in figure 4.4:

$$H_0: \text{Equal task completion time between experienced and inexperienced participants}$$
$$H_A: \text{Unequal task completion time between experienced and inexperienced participants}$$

If \bar{x}_1 equals the mean time for experienced, and \bar{x}_2 the mean time for inexperienced participants, the hypothesis can be written as:

$$H_0: \bar{x}_1 = \bar{x}_2$$
$$H_A: \bar{x}_1 \neq \bar{x}_2$$

Since we cannot assume equal variances in the two samples (Table 4.10), this test will use the Welch's t-test for unequal variances [(Walpole et al., 2012), p. 345]. Equation 4.1 is still valid. The obtained values from the test are shown in the box below. The obtained T-statistic is smaller than the critical value. The t-test therefore conclude that there is a significant difference between the means of the two population samples with a confidence interval of 95%.

Two sample, two-way t-test
Sample 1 and 2

Degree of freedom (v): 447

Significance level (α): 0.05

Critical value: 1.960

$T - \text{statistic}$: -60.442

Using equation 4.1, the absolute value of the $T - \text{statistic}$ is larger than the critical value ($|60.442| > 1.960$) and the null hypothesis is rejected and H_A accepted.

Test if experienced or inexperienced participants finish the task fastest

Because there was a statistical significant difference between time spent on each task between the participants, this section will test which group finished the task fastest. The second hypothesis tested in this section is number two in figure 4.4:

$$H_0: \text{Experienced participants do not finish the tasks faster}$$
$$H_A: \text{Experienced participants finish the tasks faster}$$

With sample 1 being experienced participants and sample 2 inexperienced participants we get the hypothesis:

$$H_0: \bar{x}_1 = \bar{x}_2$$
$$H_A: \bar{x}_1 < \bar{x}_2$$

This test gives the same T-statistics as the previous test, but the critical value is changed since this test used in the second hypothesis is a two sample, one-way t-test. The Welch's test is used since we cannot assume equal variances in the two samples. Obtained $T - \text{statistic}$ is still smaller than the critical value ($-64.654 < 1.645$). Our test is to check if the mean value of sample 1 is significantly larger than the mean value of sample 2. We use the comparison test written in equation 4.2, section 4.2.2.1. Our $T - \text{statistic}$ is not larger than the critical value, and we need to accept the null hypothesis. There is no evidence that experienced participants use less time on the tasks than the inexperienced.

4. RESULTS

Two sample, one-way t-test

Sample 1 and 2

Significance level: 5%

T – statistic: -60.442

Degree of freedom (v): 447

Significance level (α): 0.05

Critical value: 1.645

T-statistic is smaller than the critical value ($-60.442 < 1.645$) and the null hypothesis is accepted.

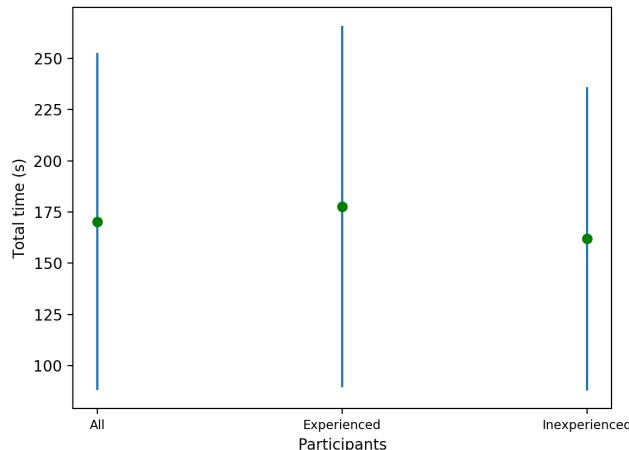


Figure 4.10: Sample 1 and 2 - mean (green dot) and standard deviation (blue line)

Since we know that there is a statistically significant difference between the two sample means, we conclude that the inexperienced participants finished the task faster than the experienced participants. The time difference can also be seen in plot 4.10. Inexperienced participants finished the tasks in average 16 seconds faster than experienced participants.

4.3.4.2 Difference between experienced and inexperienced participants in total correct elements

This section will test if there is a difference between experienced- and inexperienced participants when looking at the number of correctly chosen elements. Sample 3 and

4, table 4.2, is the correct samples to use in this test. Neither samples followed the normal distribution (table 4.9), and we need to use a non-parametric method. Both samples have ties (identical observations), and, as mentioned in section 4.2.2.3, the Mann-Whitey U test is then preferred. From histogram A.1a and A.1b we see that the samples are identical in some cases. Mann-Whitey U test should, therefore, be used to compare the population medians. The hypothesis to be tested is number three in figure 4.4

$$H_0: \text{median}_3 = \text{median}_4$$

$$H_A: \text{median}_3 \neq \text{median}_4$$

Using equation 4.5 in section 4.2.2.3 and the obtained $U - \text{statistic}$, we conclude that there is not enough evidence to reject the null hypothesis with a confidence interval of 95%. The $U - \text{statistic}$ is larger than the critical value, and the null hypothesis is accepted.

Two sample t-test

Sample 3 and 4

Significance level: 5%

$U - \text{statistic}$: 17012

Significance level (α): 0.05

Sample size, n1: 229

Sample size, n2: 200

Critical-value: 127

$U - \text{statistic}$ is larger than the critical value ($17012 > 127$)
and the null hypothesis is accepted.

4. RESULTS

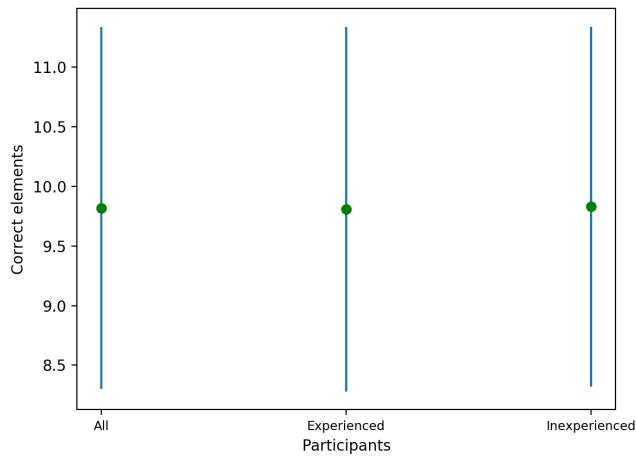


Figure 4.11: Sample 3 and 4 - mean (green dot) and standard deviation (blue line)

Results from this section show that there is not enough evidence to conclude that there is any difference between experienced and inexperienced participants when looking at the number of correctly chosen elements per task. We conclude that experienced and inexperienced participants did equally well on the task. This can also be seen visually in figure 4.11. Mean values shown in the figure is approximately equal between the participants.

4.3.4.3 Test if total time differs between the three tasks

This section will test if time spent varies between each of the three tasks (hypothesis one in figure 4.5). Sample 5, 6, and 7, table 4.3, is used in this test. The one-way *ANOVA* method will be applied to answer the hypothesis. All three samples come from populations with equal variances (4.10), the samples are also normally distributed after a Box-Cox transformation (table 4.9).

$$H_0: \bar{x}_5 = \bar{x}_6 = \bar{x}_7$$
$$H_A: \text{Total time differs between at least two of the tasks}$$

Using equation 4.4 in section 4.2.2.2 and results obtained from the calculations, the one-way *ANOVA* test rejects the null hypothesis. The obtained f -value is lower than the critical value. With a confidence interval of 95% we claim that there is a difference between the mean value of the three tasks.

One-way ANOVA
Sample 5, 6 and 7

Significance level (α): 0.05

$$v_1 = 2, v_2 = 426$$

Critical-value: 3.00

f -value: 2123.308

f -value is significantly higher than the critical value
($2123.308 > 3.00$) and the null hypothesis is rejected, H_A is
accepted

When rejecting the null hypothesis, *Tukey's method* is used to make comparisons between task A, B, and C. This test did not find any statistically significant difference between the three tasks. Visual evaluation of figure 4.12 show that task A was completed slightly faster than the two other tasks.

Tukey's test
Sample 5, 6 and 7
Significance level: 5%

Task A and Task B do not differ significantly
Task A and Task C do not differ significantly
Task B and Task C do not differ significantly

4. RESULTS

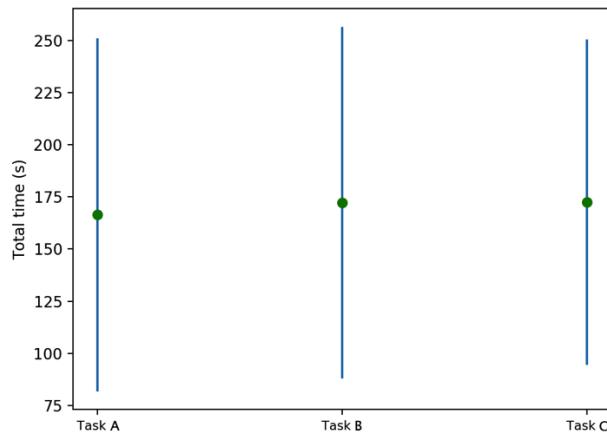


Figure 4.12: Sample 5, 6 and 7 - mean (green dot) and standard deviation (blue line)

Based on the results from this section, we conclude that there is a statistically significant difference in total time between at least two of the tasks, but the difference is not enough for *Tukey's test* to notice a difference.

4.3.4.4 Test if the number of correct elements differs between the three tasks

This section will test if there is a difference in the number of correctly chosen elements between the three tasks which are hypothesis two in figure 4.5. The test will use sample 8, 9 and 10 from table 4.4. Since sample 8 is not normally distributed (table 4.9) a non-parametric test should be applied. The Kruskal-Wallis test is the non-parametric equivalent to one-way *ANOVA* (4.2.2.4). The method tests equality of medians when the samples do not follow the normal distribution. The hypothesis tested is:

$$H_0: median_8 = median_9 = median_{10}$$

H_A : Number of correctly chosen elements differ between at least two of the tasks

Using equation 4.6 in section 4.2.2.4, the Kruskal-Wallis test rejects the null hypothesis. The obtained H -value is smaller than the critical value. The p-value is approximately zero, and this gives a good indication that the result is significant. With a confidence interval of 95%, we claim that there is a difference between the median value of the three tasks.

Kruskal-Wallis test
Sample 8, 9 and 10

Significance level (α): 0.05

$$v = 2$$

Critical-value: 5.991

$P - value$: $3.967 * 10^{-72}$

$H - value$: 328.816

$H - value$ is significantly higher than the critical value
($328.816 > 5.991$) and the null hypothesis is rejected, H_A is accepted

Like in the one-way ANOVA, a *post hoc* test should be used to make paired comparisons to determine which groups differ. The *post hoc* test applied is Tukey's test. Results from Tukey's test resulted in a significant difference in the number of correctly chosen elements between task A and task B, and task A and task C. This can also visually be seen in figure 4.12. The participants had in average 0.8 more correct elements in task A compared to the two other tasks. Task A also has a smaller standard deviation than the other tasks.

Tukey's test
Sample 8, 9 and 10
Significance level: 5%

Task A and Task B differs significantly

Task A and Task C differs significantly

Task B and Task C does not differ significantly

4. RESULTS

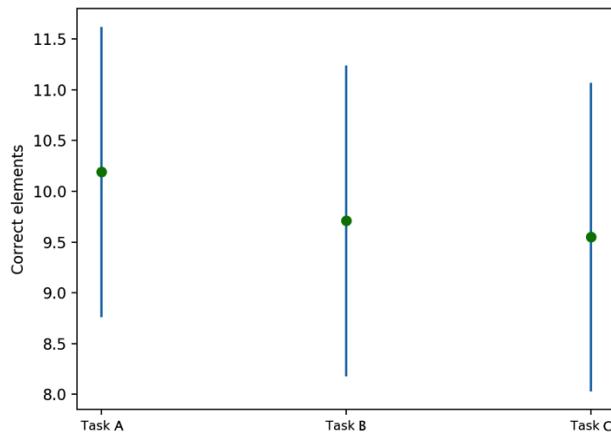


Figure 4.13: Sample 8, 9 and 10 - mean (green dot) and standard deviation (blue line)

The results in this section find statistically significant evidence that the participants had a higher number of correct elements in task A than in the two other tasks.

4.3.4.5 Test differences in results from experienced participants

This section will answer two hypotheses about experienced participant's results divided by task. The first hypothesis will test if time spent on each of the three tasks differ and the second hypothesis will test if the number of correct elements in each of the three tasks differs when the samples only include results from experienced participants. Sample 11, 12 and 13, from table 4.5, will be used to answer the first hypothesis. All three samples are normally distributed after a Box-Cox power transformation (4.9) and also come from populations with equal variances (4.10). Sample 14, 15 and 16, from table 4.6, will be used on the second hypothesis. These three samples are also normally distributed (4.9) and come from populations with equal variances (4.10). The one-way ANOVA method will be used to test both hypotheses.

The first hypothesis is:

$$H_0: \bar{x}_{11} = \bar{x}_{12} = \bar{x}_{13}$$
$$H_A: \text{Total time differs between at least two of the tasks}$$

Using equation 4.4 from section 4.2.2.2, the one-way ANOVA test rejects the null hypothesis. The obtained *f-value* (1216.919) is higher than the critical value (3.00). The calculated p-value is approximately zero, which gives a good indication that the

result is significant. With a confidence interval of 95% the author claim that there is a time difference between the three tasks.

Tukey's *post hoc* test did not find any significant difference between the three tasks with a significance level of 5%. Figure 4.14 show an approximately similar mean value in all three tasks. Task B has a lower mean and less standard deviation, but it is not statistically significantly different.

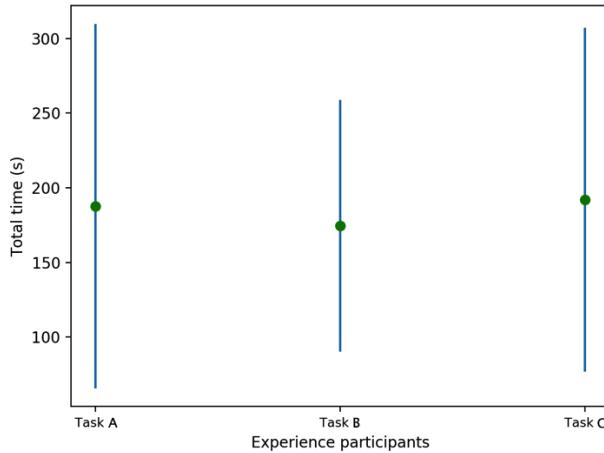


Figure 4.14: Sample 11, 12 and 13 - mean (green dot) and standard deviation (blue line)

The second hypothesis is:

$$H_0: \bar{x}_{11} = \bar{x}_{12} = \bar{x}_{13}$$

H_A : Number of correct elements in each task differs between at least two of the tasks

Using equation 4.4, the one-way ANOVA test rejects the null hypothesis. The obtained *f-value* (8.210) is higher than the critical value (3.00). The p-value is approximately zero, which gives a good indication that the result is significant. With a confidence interval of 95%, we claim that there is a difference between the mean value of at least two of the tasks.

Tukey's *post-hoc* test resulted in a significant difference in the number of correctly chosen elements between task A and task B, and task A and task C. Figure 4.15 show that task A has a higher mean value than the two other tasks. Task A also has a smaller standard deviation.

4. RESULTS

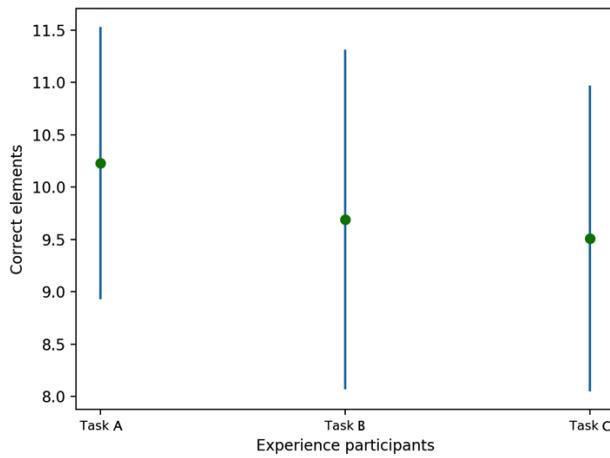


Figure 4.15: Sample 14, 15 and 16 - mean (green dot) and standard deviation (blue line)

With the results found in this section, we conclude that there is a statistically significant difference in the number of correctly chosen elements between the three tasks for experienced participants. They got the best result on task A. Time spent on each task also differs between at least two of the tasks, but the difference is not significant enough so that Tukey's test can determine a difference. Figure 4.14 show that task B has the lowest mean time value of the three tasks.

4.3.4.6 Test differences in results from inexperienced participants

This section will answer the same hypothesis as the previous section, but using results from inexperienced participants. The first hypothesis will test if time spent on each task differ and the second hypothesis will test if the number of correct elements in each task differs. Sample 17, 18 and 19, from table 4.7, will be applied in the first hypothesis test. These samples are normally distributed (table 4.9) and come from populations with equal variances (4.10). Sample 20, 21 and 22, from table 4.8, will be used on the second hypothesis. All three samples are normally distributed (table 4.9) and also come from populations with equal variances (4.10). The one-way ANOVA will be used to test both hypotheses.

The first hypothesis is:

$$H_0: \bar{x}_{17} = \bar{x}_{18} = \bar{x}_{19}$$
$$H_A: \text{Total time differ between at least two of the tasks}$$

The one-way ANOVA test rejects the null hypothesis and accepts the alternative hy-

pothesis (H_A) with a significance level of 5%. The obtained f-value from the test is higher than the critical value ($905.34 > 3.00$). Since the alternative hypothesis was accepted, Tukey's *post hoc* test is used to make compared comparisons between task A, task B and task C. The test does not find a significant difference when comparing each of the three tasks with a significance level of 5%. Figure 4.16 show that inexperienced participants spent more time on task B than the other tasks, but the difference is not significant according to Tukey's test.

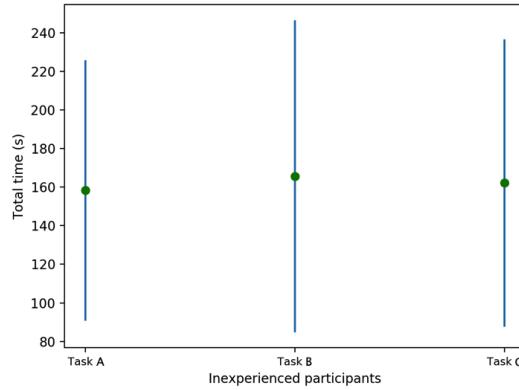


Figure 4.16: Mean (green dot) and standard deviation (blue line) for sample 17, 18 and 19

The second hypothesis is:

$$H_0: \bar{x}_{20} = \bar{x}_{21} = \bar{x}_{22}$$

H_A : Number of correct elements differ between at least two of the tasks

The one-way ANOVA test rejects the null hypothesis (H_0) and accepts the alternative hypothesis (H_A) with a significance level of 5%. The obtained f-value from the test is higher than the critical value ($189.05 > 3.00$). Tukey's *post hoc* test cannot find a significant difference when comparing the tasks with a significance level of 5%. Looking at figure 4.17, task A has a higher mean value than the two other tasks, and task B also has a higher mean than task C. Even though there are differences in the number of correct elements between the tasks, it is not significant according to Tukey's test.

4. RESULTS

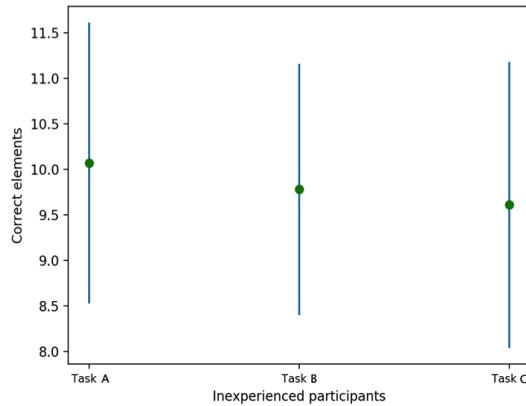


Figure 4.17: Mean (green dot) and standard deviation (blue line) for sample 20, 21 and 22

With the results found in this section, we conclude that there is a statistically significant difference in both total time spent on each task and the number of correctly chosen elements between at least two of the tasks. The differences are not significant enough so that Tukey's test can determine which task differs. Figure 4.16 and 4.17 show that task A has the lowest mean time value and the highest mean correct value.

4.3.4.7 Hypothesis test summary

Table 4.11: Summary of hypothesis tests computed in section 4.3.4

Hypothesis (Dependent variable, Independent variable)	Participants Sample ID	Hypothesis is accepted
Total time, Experienced, Inexperienced There is a difference between experienced and inexperienced participants Experienced participants finish the tasks faster than inexperienced Inexperienced participants finish the tasks faster than experienced	All 1 and 2 1 and 2 1 and 2	Yes No Yes
Correct elements, Experienced, Inexperienced There is a difference between experienced and inexperienced participants	All 3 and 4	No
Total time, Task A, Task B, Task C Total time is different between at least two of the tasks Task A significantly differs from Task B Task A significantly differs from Task C Task B significantly differs from Task C	All 5, 6 and 7 5 and 6 5 and 7 6 and 7	Yes No No No
Correct elements, Task A, Task B, Task C The number of correctly chosen elements is different between at least two of the tasks Task A significantly differs from Task B Task A significantly differs from Task C Task B significantly differs from Task C	All 8, 9 and 10 8 and 9 8 and 10 9 and 10	Yes Yes Yes No
Total time, Task A, Task B, Task C Total time is different between at least two of the tasks Task A significantly differs from Task B Task A significantly differs from Task C Task B significantly differs from Task C	Experienced 11, 12 and 13 11 and 12 11 and 13 12 and 13	Yes No No No
Correct elements, Task A, Task B, Task C Number of correct elements differs between at least two of the tasks Task A significantly differs from Task B Task A significantly differs from Task C Task B significantly differs from Task C	Experienced 14, 15 and 16 14 and 15 14 and 16 15 and 16	Yes Yes Yes No
Total time, Task A, Task B, Task C Total time is different between at least two of the tasks Task A significantly differs from Task B Task A significantly differs from Task C Task B significantly differs from Task C	Inexperienced 17, 18 and 19 17 and 18 17 and 19 18 and 19	Yes No No No
Correct elements, Task A, Task B, Task C	Inexperienced	

4. RESULTS

Number of correct elements differs between at least two of the tasks	20, 21 and 22	Yes
Task A significantly differs from Task B	20 and 21	No
Task A significantly differs from Task C	20 and 22	No
Task B significantly differs from Task C	21 and 22	No

5 | Discussion

Our results have found support for the possibility of giving geospatial micro-tasks to all individuals, independent of background. There was statistically significant evidence that inexperienced participants finished the tasks faster than experienced participants, this is also shown in figure 4.10. The mean time difference was 16 seconds, a statistically significant, but relatively small difference. Gadiraju et al. (2015b) demonstrated that using a proper training section decreased the task completion time with up to 41%. An assumption is that inexperienced participants followed the instruction video and training task more carefully than experienced participants. One would expect experienced participants to finish the tasks faster since they are familiar with map interaction and interpreting base maps and meta information. The training section was maybe enough to eliminate the inexperienced uncertainty towards the tasks, making them just as fast as the experienced. The author also suspects the experts not to use the training section thoroughly, resulting in them completing the tasks slower than the non-experts.

During the pilot test, the author noticed that the experienced participants used the map aids given in question one more frequently than inexperienced. The map aids were zoom, panning and a layer control to show/hide the building footprints. Without the aids the participants could not see the shape of the building on the base map, so they were important to use. A possible explanation for the time difference is that the experienced participants spent more time using the map aids provided since they are more familiar with them and knew how to use them. Using the aids more frequently to investigate the base map and layers more thoroughly to be more confident in their answers. It is a possibility of them executing the tasks more complicated than necessary. The experienced individuals have most probably solved much more complex tasks and maybe compares the micro-tasks with similar and more advanced tasks. They then overestimate the complexity resulting in them spending more time on each micro-task.

When analyzing the number of correctly chosen elements between experienced and inexperienced participants, the difference is not statistically significant. Figure 4.11 show a very similar mean value between them. Inexperienced participants had a mean of 9.83, while experienced had a mean of 9.81 correct elements. One would expect that inexperienced participants had fewer correct elements since they spent less time on the tasks. Salk et al. (2016) results also observed minor differences between the professional and non-professional participants. They concluded that professional background had a limited first-order relationship with task accuracy. It can be argued that the introduction video and training task remove the effect of having a professional background. Training made the micro-tasks understandable for everyone so that professional background had no effect on the quality of the task results. Gadiraju et al. (2015b) showed that with a training section the task quality increased up to 5%. A 5% quality increase can be enough to remove the differences in task performance between experts and non-experts. See et al. (2013) also concluded that with proper targeted

5. DISCUSSION

training material the differences between experts and non-experts could decrease. A conclusion that matches this thesis results well.

The splitting of participants into experienced and inexperienced was based on the question "Do you have experience of working with geospatial data?". Other studies ask the participants for background information through a registration procedure. See et al. (2013) and Salk et al. (2016) considered people with a background in remote sensing/spatial science as experts, and people who were new to the discipline or had a self-declared limited background as non-experts. In this study, the participants are self-declared experts / non-experts. It is not possible to validate this information. In the pilot-test we knew the background of the participants, they answered yes and no on the experience question as the author anticipated. With this observation, we conclude that overall the participants correctly declared themselves as experienced / inexperienced.

There were minor differences between the three tasks. Results found no statistical support for faster completion time with fewer elements. Time spent completing each task was approximately the same. Figure 4.14 show a slightly faster task completion on task A, but not a statistically significant difference according to *Tukey's test*. Experienced participants finished task B fastest (figure 4.14) and inexperienced finished task A fastest (figure 4.16) but this is also not statistically significant. The small differences in time show that participants found all three tasks manageable to solve. Fewer elements in the micro-tasks did not reduce the task completion time.

One can question the time variable because it does not reflect how much time the participant spent in front of the screen from task start to task end. It reflects how much time passed when solving the two questions or micro-tasks. It can be argued that the participants spend more time on task A in total. Time spent switching to the next question and fetching the next task element is not added to the time variable. In total the participant probably spent more time in front of the screen doing task A compared to task C. In task C all six task elements are present, so the participants did not have to wait for the web application too, i.e., fetch the next task element. The Los Angeles building import used another strategy. Their approach can be described as a combination of task A and task C. The buildings were imported one by one, but in their solution, all buildings covering a selected area was visible on the map at the time time. The map window highlighted one and one building, guiding the workers through every building. This made it easier for the micro-tasking workers to concentrate on each building and at the same time have all elements available. This approach eliminates time spent fetching and switching to the next building footprint on the map, but at the same time probably keeps the task performance high. Future studies should look closer on this task approach.

Looking at the quality of the task results, task A is statistical significantly better than the two other tasks according to *Tukey's test*. Our results have found support for the quality to increase with fewer elements present in the task, also shown in figure 4.13. Using the figure and table 4.4, participants had in average one more correct element in task A compared to the two other tasks. Participants did worse on task C, but the

difference is small. Experienced participants got better results on task A, and this was also statistically significant. Inexperienced participants also got the best results on task A, but this is not statistically significant. The quality of task results was more even in the three tasks for inexperienced than experienced participants. This can be seen in figure 4.15 and 4.17.

Task A had an average difficulty rating of 2.11, task B of 2.14 and task C of 2.5. Experienced participants gave a lower difficulty score than inexperienced on all three tasks (in average a 0.20 lower difficulty score). All scores are surprisingly low. In the "after each survey" form, score one was described as easy and score five as hard. It is clear that the participants overall found the tasks manageable to solve.

One can question how realistic the experiment application is compared to a real micro-tasking situation. The main difference is the absence of motivation factors in the experiment. The author decided to neglect these factors. Gamification and payments are frequently used motivation factors, as shown in chapter 2. There is likely a difference between a curious participant joining the experiment and, i.e., an MTurk worker who gets paid per completed micro-task. To achieve a perfectly realistic study the micro-tasks used in the experiment could be posted on a micro-tasking platform, i.e., MTurk or CrowdFlower. Using a micro-tasking platform would require funds to pay the workers, which is something this master thesis did not have. Another strategy could be to create a game-like approach as Luis von Ahn did through, i.e., "The ESP game" (section 2.2). Solving the micro-tasks though a game would exploit the motivation factor. A game would probably be challenging to distribute and get people to invest time playing it. A new, unknown game is not easy to distribute and would also demand much time and effort to develop.

Much of the work the OpenStreetMap community has made around micro tasks so far has been correct. They are fixing geometry errors through a gamification approach using micro-tasking. They also divide large import tasks into smaller tasks distributing them on their micro-tasking platform (section 2.4.1.2). What the author evaluate as missing knowledge in OSM, is partitioning the large tasks. This was a problem OSM encountered during the New York building import. The import dataset was divided into predefined regions, containing different building density. This gave a huge deviation between the number of elements in each micro-task. The tasks could not be distributed to all workers on the Tasking Manager platform, because of the time and resources spent training new workers (shown in section 2.4.4). One can question that each micro-tasks incorporated too many buildings, making the tasks too complex. Neither of the three tasks given in this experiment was considered as too complex by the participants when evaluating the difficulty score. The author would recommend dividing a large task by a predefined number and not by predefined regions as both import jobs mentioned in section 2.4.4 did.

6 | Conclusion

This thesis amplifies the assumption that micro-tasking works well together with geospatial data. We conclude that it is not necessary to only use experienced individuals when solving geospatial micro-tasks, especially if the platform publishing the tasks are well designed and provides a training session. Having the ability to use all humans, independent of background, makes it easier to reach a considerable larger mass of people to solve the micro-tasks. We also see an indication that the quality of the task results will increase with fewer elements in each micro-task. Time spent to complete the micro-tasks is not significantly affected by the number of elements, but the time spent in total will probably increase with fewer elements. Therefore, the method used by Los Angeles building import can be another approach to how elements in micro-tasks can be managed. This method needs further research.

The results from this thesis can be used as guidance on how to break down a large geospatial task. If quality is the key factor, the large task must be divided so that the micro-tasks contains few elements. The downside of having too few elements in the micro-tasks is the time spent on not task specific operations. If the worker needs to click and wait for the next element to load, the workers will spend unnecessary time. If quality mechanisms already are implemented it is possible to break the task into larger chunks (containing, i.e., six elements). The number of mistakes will probably increase with the number of elements present in the micro-tasks, but if the quality mechanisms are well developed, the errors will be handled. Benefits of dividing each task into larger chunks are that the worker will most likely spend less time in total doing the micro-tasks. We would not recommend chunks containing an uneven number of elements, like in the New York building import.

The author would recommend organizations to exploit micro-tasking when it is possible to break the large geospatial task into smaller parts. It is shown that the method makes imports of large datasets in OpenStreetMap easier. The method is also exploited in machine-learning, combining the human computation with the power of computers. Humans manage to evaluate and validate geometric layers, shown through question one and also interpret meta information, shown through question two in the experiment. The theory of which geospatial tasks that are most suitable to solve though micro-tasking is not known yet. This thesis indicates that task design and training sections are almost as important as the content of the tasks.

Appendices

A | Normal tests

A.1 All participants, number of correctly chosen elements variable

This section will test if sample 3 and 4 (table 4.2) are drawn from a normal distribution. These samples will be used to test if there are any difference in the number of correctly chosen elements per task between experienced and inexperienced participants.

A visual analysis of the samples histogram A.1a and A.1b, we see a good indication that sample 3 and 4 are not drawn from a normal distribution. Both are clearly negatively skewed.

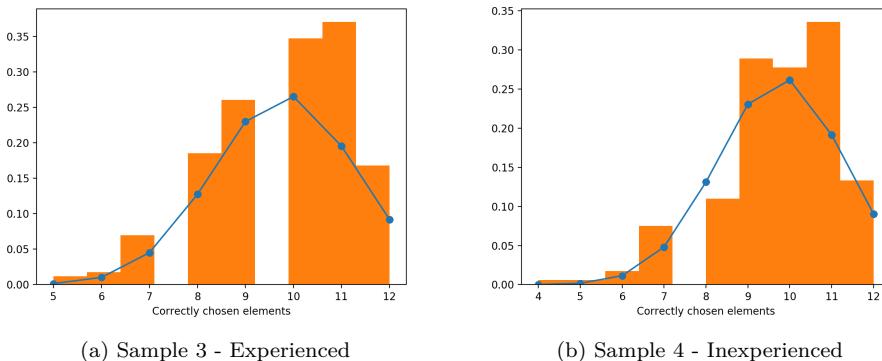


Figure A.1: Histograms with normal distribution fit with samples containing the number of correctly chosen elements

D'Agostino and Pearson normality test confirm our visual analysis. Both samples accept the alternative hypothesis with p-values (0.00443, 0.00013) lower than the significance level (0.05). The null hypothesis is rejected and H_A accepted for sample 3 and 4.

Sample 3 and 4 is Box-Cox power transformed because the null hypothesis was rejected. After the transformation, a new D'Agostino and Pearson normality test was performed. Both samples also failed this test and the alternative hypothesis (H_A) is accepted. Hypothesis including sample 3 and 4 need to be tested with non-parametric methods.

A.2 All participants divided by task, total time variable

In this section sample 5, 6, and 7 (table 4.3) is tested for a normal distribution. These samples will be used to test whether there is a significant difference between the three tasks when considering the total time variable.

A visual analysis of the three histograms in figure A.2a, A.2b and A.2c show a positive skewness, just like the histograms in figure 4.6. This gives an indication that the three samples do not follow the normal distribution.

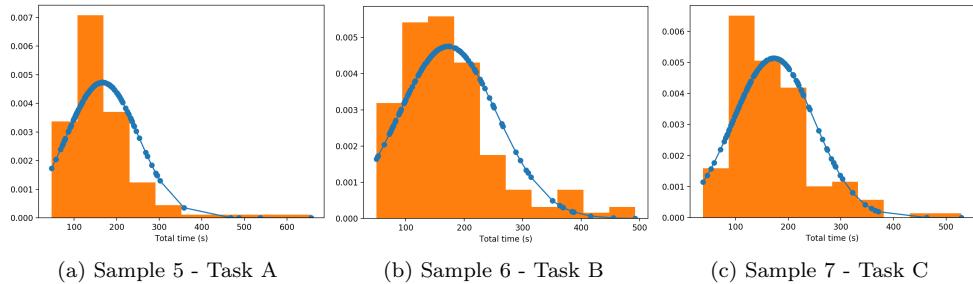


Figure A.2: Histogram with normal distribution fit - sample with total time per task

The D'Agostino and Pearson normality test agreed with the visual analysis. Obtained p-values for all three samples ($2.39 * 10^{-24}$, $2.57 * 10^{-9}$, and $1.71 * 10^{-11}$) are smaller than the significance level (0.05), and the null hypothesis is rejected. The samples do not follow the normal distribution with a confidence interval of 95%.

Because the null hypothesis was rejected, the samples are Box-Cox power transformed. Histograms of each sample after the transformation is shown in figure A.3a, A.3b and A.3c. A visual analysis of the histograms gives a good indication that the transformed data is approximately normally distributed. The histograms have a skewness of approximately zero.

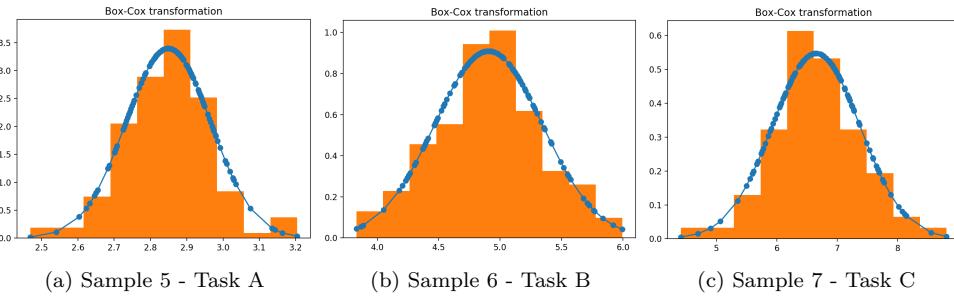


Figure A.3: Histogram with normal distribution fit after Box-Cox transformation, total time variable

The D'Agostino and Pearson normality test confirms the visual interpretation. The p-values (0.164, 0.982, and 0.354) of all three samples are higher than the significance level (0.05), and the null hypothesis is accepted. Sample 5, 6 and 7 follows the normal distribution after the transformation, and parametric methods can be used with these samples.

A.3 All participants divided by task, correct element variable

This section will examine sample 8, 9 and 10 (table 4.4) for the normal distribution assumption. These samples will be used to test whether there is a significant difference between the three tasks when looking at the number of correctly chosen elements variable.

A visual analysis of the three histograms in figure A.4a, A.4b and A.4c show a negative skewness. This gives an indication that the three samples are not drawn from a normal distribution.

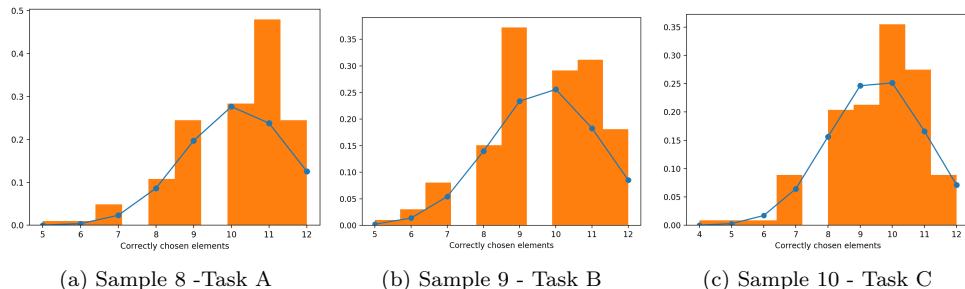


Figure A.4: Histogram with normal distribution fit showing samples with number of correct elements per task

D'Agostino and Pearson normality test confirms our visual analysis of the histograms in two of three samples. Sample 9 passes the normality test, even though the p-value (0.099) is close to the significance level (0.05). Sample 8 and sample 10 do not pass the normal assumption test. Both samples obtained a p-value (0.00022 and 0.0047) smaller than the significance level. The null hypothesis is rejected for sample 8 and 10, and the alternative hypothesis is accepted. The null hypothesis is accepted for sample 9.

A Box-Cox power transformation is applied to all three samples. A transformation changes the data, and to correctly compare the results sample 9 has to be transformed, even though it follows a normal distribution. The transformed data is shown in histogram A.5a, A.5b, and A.5c. All three are negatively skewed, sample 9, and 10 less than sample 8. The conclusion is not obvious in these histograms.

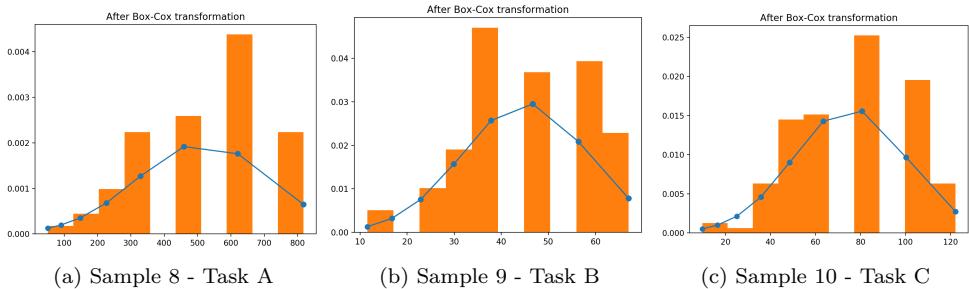


Figure A.5: Histogram with normal distribution fit after Box-Cox

D'Agostino and Pearson normality test accepts the null hypothesis on sample 9 and 10 and rejects it on sample 8. Sample 9 and 10 has p-values (0.0752 and 0.2104) higher than the significance level, while computed p-value for sample 8 (0.0027) is significantly lower. When using these three samples in hypothesis tests, a non-parametric method should be used. This is because sample 8 do not follow the normal distribution.

A.4 Experienced participants divided by task, total time variable

In this section, sample 11, 12, and 13, shown in table 4.5, is tested if they follow a normal distribution. The three samples will be used to test whether there is a significant difference between the total time results for the three tasks when considering only experienced participants.

A visual interpretation of the histograms in figure A.6 show that all three samples are positively skewed (figure 4.2). Skew gives a fairly strong evidence that the samples are not normally distributed.

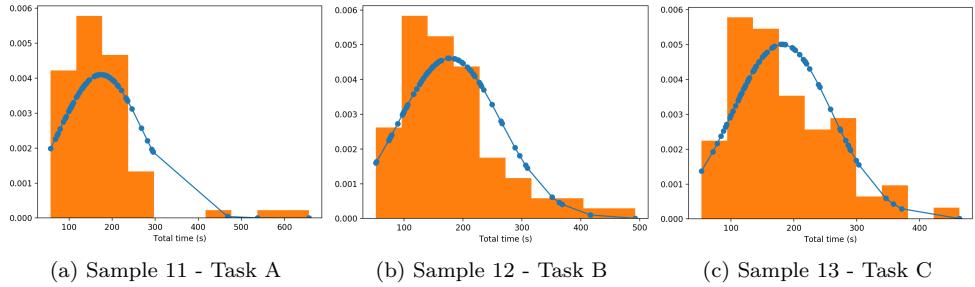


Figure A.6: Histograms with normal distribution fit with samples containing total time to complete each task

D'Agostino and Pearson normality test confirms our visual interpretation of the three histograms. All three p-values ($1.229 * 10^{-14}$, $2.678 * 10^{-5}$ and 0.000884) are lower than the significance level (5%). Sample 11, 12 and 13 do not pass the normality assumption with a confidence interval of 95%.

A Box-Cox power transformation is applied to all three samples since the null hypothesis was rejected. Histograms with normal distribution fit containing the transformed data is shown in figure A.7. Visually, the histograms look like they follow a normal distribution with minimal skewness.

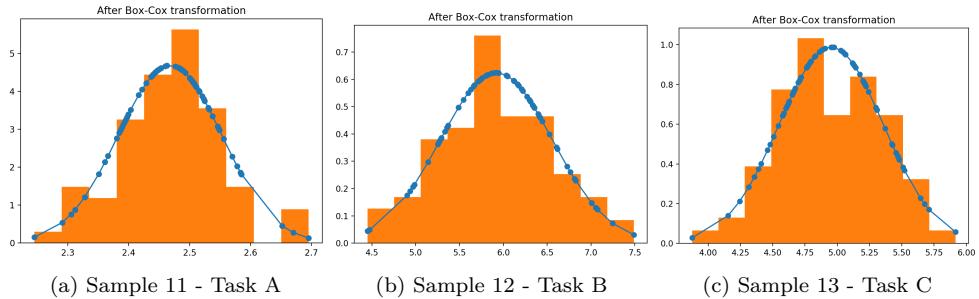


Figure A.7: Histograms with normal distribution fit containing Box-Cox transformed data

The Box-Cox transformed data is tested with D'Agostino and Pearson normality test. All three samples obtained p-values (0.694, 0.955 and 0.887) larger than the significance level (0.05). Within a confidence interval of 95%, the test concludes that sample 11, 12 and 13 is normally distributed. These samples can be used in parametric methods.

A.5 Experienced participants divided by task, correct elements variable

This section will test if sample 14, 15, and 16, shown in table 4.6, follows the normal distribution. The three samples will be used to test whether there is a significant difference in the number of correct elements between the three tasks when comparing only experienced participants.

A visual interpretation of the histograms in A.8 show that all three samples are slightly negatively skewed. Sample 15 (A.8a) and sample 16 (A.8b) has less skewness than sample 14 (A.8c).

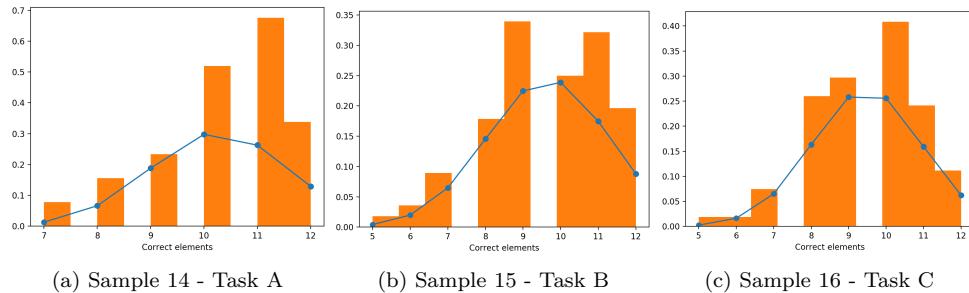


Figure A.8: Histogram with normal distribution fit showing samples with number of correct elements results for experienced participants

D'Agostino and Pearson normality test confirms our visual interpretation of the three histograms. All three p-values (0.0588, 0.2067 and 0.2975) are higher than the significance level (0.05). Notice that sample 14 has a lower p-value than the two other samples. This sample is not as significant as the two other samples. Sample 14, 15 and 16 pass the normality test with a confidence interval of 95%. The null hypothesis is accepted. These samples can be used in parametric methods.

A.6 Inexperienced participants divided by task, total time variable

This section will test if sample 17, 18 and 19, table 4.7, follows the normal distribution. These samples will be used to test whether there is a significant difference in total time between the three tasks when only looking at inexperienced participants.

A visual analysis of the histograms A.9a, A.9b and A.9c show a positive skewness. The skew is less than the histograms in figure A.6, but is most likely too large for the samples to be normally distributed.

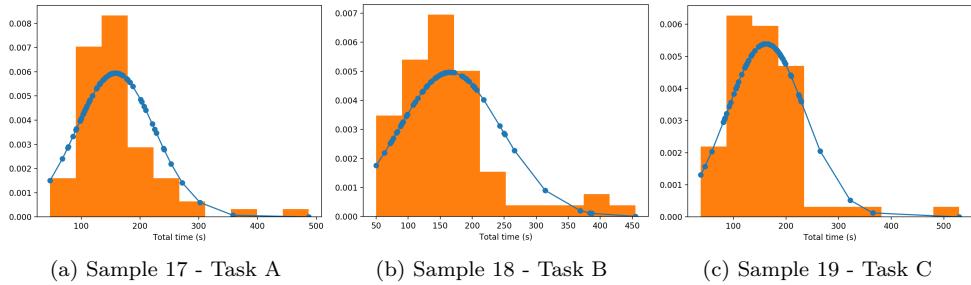


Figure A.9: Histogram with normal distribution fit

The D'Agostino and Pearson normality test agrees with the visual analysis. The three obtained p-values ($1.586 * 10^{-11}$, $1.773 * 10^{-6}$ and $2.312 * 10^{-11}$) are all significantly lower than the significance level of 0.05. Sample 17, 18, and 19 do not follow a normal distribution with a confidence interval of 95%. The null hypothesis is rejected.

The samples are Box-Cox power transformed. The histograms after transformation (A.10a, A.10b, and A.10c) are visually evaluated to be normally distributed.

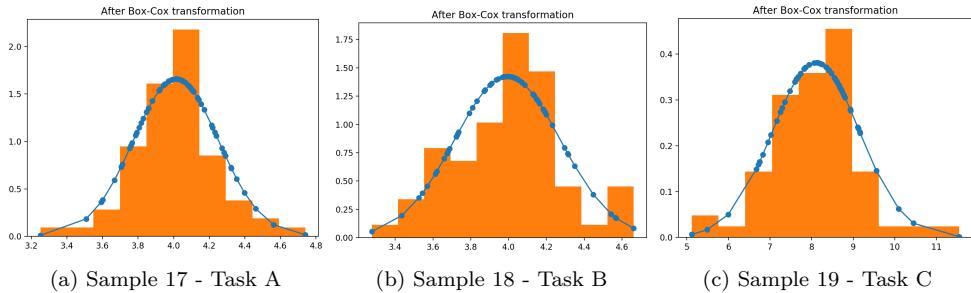


Figure A.10: Histogram with normal distribution fit after Box-Cox

A new D'Agostino and Pearson normality test on the transformed data confirms that all three samples are drawn from a normal distribution with a significance level of 5%. The obtained p-values (0.139, 0.909 and 0.067) are larger than 0.05, and the null hypothesis is accepted. These samples can be used in parametric methods

A.7 Inexperienced participants divided by task, correct elements variable

This section will test if sample 20, 21, and 22, in table 4.8, follows the normal distribution. These samples contain the number of correctly chosen elements in each of the

APPENDIX A. NORMAL TESTS

three tasks from only inexperienced participants. The samples will be used to test if inexperienced participants do better in one of the tasks.

A visual interpretation of histogram A.11a, A.11b and A.11c show a negative skew, similar to the histograms containing results from only experienced participants (A.8).

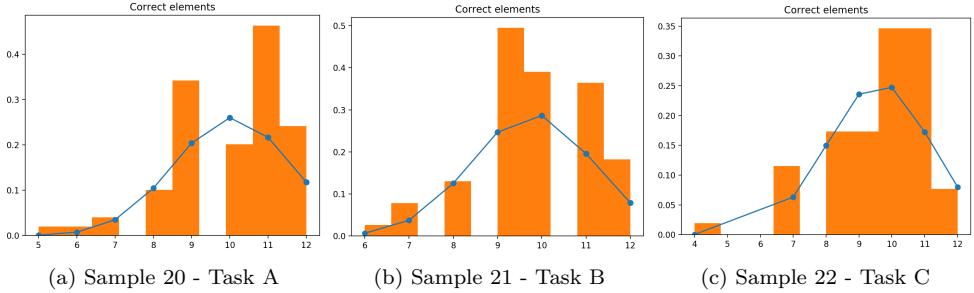


Figure A.11: Histogram with normal distribution fit

The D'Agostino and Pearson normality test rejects the null hypothesis on sample 20 and 22 and accepts the null hypothesis on sample 21. Sample 20 and 22 obtained p-values (0.007 and 0.004) lower than 0.05 and sample 21 obtained a p-value (0.523) higher than 0.05. The test concludes that sample 21 follows the normal distribution, and sample 20 and 22 does not with a significant level of 5%.

A Box-Cox power transformation is applied to all three samples. The transformation changes the data, and to correctly compare the data, sample 21 also has to be transformed, even though the original data was followed the normal distribution. The transformed samples are shown in histogram A.12a, A.12b and A.12c. All three histograms are less skewed than the original histograms (A.11).

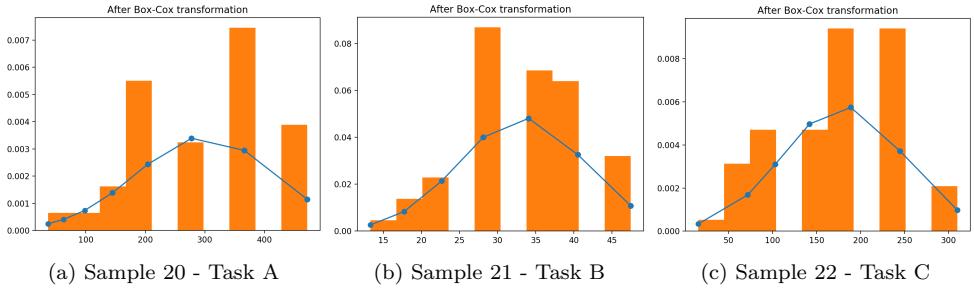


Figure A.12: Histogram with normal distribution fit after Box-Cox transformation

The D'Agostino and Pearson normality method is executed on the transformed samples. The three obtained p-values (0.061, 0.714 and 0.311) are higher than 0.05. The null hypothesis is accepted. Sample 20, 21 and 22 follows the normal distribution with a significance level of 5% and can be used in parametric methods.

Bibliography

- Affairs, A. S. f. P. (2013). System Usability Scale (SUS).
- Bangor, A., Kortum, P., and Miller, J. (2009). Determining What Individual SUS Scores Mean: Adding an Adjective Rating Scale. *Journal of Usability Studies*, 4(3):114–123.
- Barrouillet, P., Bernardin, S., Portrat, S., Vergauwe, E., and Camos, V. (2007). Time and Cognitive Load in Working Memory.
- Barth, A. (2014a). OpenStreetMap | lxbarth sin dagbok | Importing 1 million New York City buildings and addresses. Date Accessed: 2017-05-23 URL: openstreetmap.org/user/lxbarth/diary/23588.
- Barth, A. (2014b). Over 1 million New York City buildings and addresses imported to OpenStreetMap | Mapbox. Date Accessed: 2017-05-23 URL: www.mapbox.com/blog/nyc-buildings-openstreetmap/.
- Ben, S. and Plaisant, C. (2009). *Designing the User Interface*. Pearson, fifth edition.
- Bernstein, M. S., Little, G., Miller, R. C., Hartmann, B., Ackerman, M. S., Karger, D. R., Crowell, D., and Panovich, K. (2015). Soylent: A Word Processor with a Crowd Inside. *COMMUNICATIONS OF THE ACM*, 58(8):85–94.
- Biewald, L. (2015). Why human-in-the-loop computing is the future of machine learning | Computerworld. Date Accessed: 2017-05-14 URL: www.computerworld.com/article/3004013/robotics/why-human-in-the-loop-computing-is-the-future-of-machine-learning.html.
- Brooke, J. (1996). *SUS-A quick and dirty usability scale*. "Usability Evaluation In Industry". Taylor & Francis.
- Chilton, S., Building, F., and Burroughs, T. (2009). CROWDSOURCING IS RADICALLY CHANGING THE GEODATA LANDSCAPE : CASE STUDY OF OPENSTREETMAP.
- Cohen, S. (2013). The Relationship of Man and Machine. Date Accessed: 2017-06-03 URL: ecorner.stanford.edu/videos/3093/The-Relationship-of-Man-and-Machine.
- community, O. (2017a). Import/Guidelines – OpenStreetMap Wiki. Date Accessed: 2017-05-23 URL: wiki.openstreetmap.org/w/index.php?title=Import/Guidelines&oldid=1447838.
- community, O. (2017b). Los Angeles, California/Buildings Import – OpenStreetMap Wiki. Date Accessed: 2017-05-23 URL: wiki.openstreetmap.org/w/index.php?title=Los_Angeles,_California/Buildings_Import&oldid=1447838.

BIBLIOGRAPHY

- contributors, O. (2017). lacounty:bld_id | Keys | OpenStreetMap Taginfo. Date Accessed: 2017-05-31 URL: taginfo.openstreetmap.org/keys/lacounty:bld_id#overview.
- Deng, X., Joshi, K. D., and Galliers, R. D. (2016). THE DUALITY OF EMPOWERMENT AND MARGINALIZATION IN MICROTASK CROWDSOURCING: GIVING VOICE TO THE LESS POWERFUL THROUGH VALUE SENSITIVE DESIGN 1. 40(2):279–300.
- Difallah, D. E., Catasta, M., Demartini, G., Ipeirotis, P. G., and Cudré-Mauroux, P. (2015). The Dynamics of Micro-Task Crowdsourcing. *Proceedings of the 24th International Conference on World Wide Web - WWW '15*, pages 238–247.
- Difallah, D. E., Demartini, G., and Cudré-Mauroux, P. (2016). Scheduling Human Intelligence Tasks in Multi-Tenant Crowd-Powered Systems. *Proceedings of the 25th International Conference on World Wide Web - WWW '16*, pages 855–865.
- Erichsen, A. S. S. (2016). Evaluation of the Micro-Tasking Method for Importing High-Detail Building Models to OpenStreetMap. github.com/annesofie/Prosjektoppgave/blob/master/openstreetmap.pdf.
- EYeka (2015). The State of crowdsourcing 2015 - How the world's biggest brands and companies are opening up to consumer creativity. Technical report.
- Fan, H., Zipf, A., Fu, Q., and Neis, P. (2014). Quality assessment for building footprints data on OpenStreetMap. *International Journal of Geographical Information Science*, (28:4):700–719.
- Franklin, M. J., Kossmann, D., Kraska, T., Ramesh, S., and Xin, R. (2011). CrowdDB: answering queries with crowdsourcing. *Proceedings of the 2011 international conference on Management of data - SIGMOD '11*, pages 61–72.
- Frost, J. (2015). Choosing Between a Nonparametric Test and a Parametric Test. Date accessed: 2017-04-23 URL: blog.minitab.com/blog/adventures-in-statistics-2/choosing-between-a-nonparametric-test-and-a-parametric-test.
- Gadiraju, U., Demartini, G., Kawase, R., and Dietze, S. (2015a). Human Beyond the Machine: Challenges and Opportunities of Microtask Crowdsourcing. *IEEE Intelligent Systems*, 30(4):81–85.
- Gadiraju, U., Fetahu, B., and Kawase, R. (2015b). Training Workers for Improving Performance in Crowdsourcing Microtasks. pages 100–114. Springer, Cham.
- Ghasemi, A. and Zahediasl, S. (2012). Normality tests for statistical analysis: a guide for non-statisticians. *International journal of endocrinology and metabolism*, 10(2):486–9.
- Gutzwiller, L. (2017). Why You Can't Remove Humans from AI Training Loops | Mighty AI. Date Accessed: 2017-05-23 URL: mty.ai/blog/why-you-can't-remove-humans-from-ai-training-loops/.

-
- Ha, A. (2016). CrowdFlower raises \$10M to combine artificial intelligence with crowdsourced labor | TechCrunch. Date Accessed: 2017-05-08 URL: techcrunch.com/2016/06/07/crowdflower-series-d/.
- Holzinger, A. (2013). Human–Computer Interaction and Knowledge Discovery (HCI-KDD): What Is the Benefit of Bringing Those Two Fields to Work Together? *Springer Lecture Notes in Computer Science LNCS 8127*, pages 319–328.
- Holzinger, A. (2016). Interactive machine learning for health informatics: when do we need the human-in-the-loop? *Brain Informatics*, 3(2):119–131.
- Holzinger, A., Plass, M., Holzinger, K., Crișan, G. C., Pintea, C.-M., and Palade, V. (2016). Towards interactive Machine Learning (iML): Applying Ant Colony Algorithms to Solve the Traveling Salesman Problem with the Human-in-the-Loop Approach. In *Availability, Reliability, and Security in Information Systems*, pages 81–95. Springer, Cham.
- Howe, J. (2006). The Rise of Crowdsourcing. *Wired Magazine*, 14(06):1–5.
- Huotari, K. and Hamari, J. (2017). A definition for gamification: anchoring gamification in the service marketing literature. *Electronic Markets*, 27(1):21–31.
- Ipeirotis, P. G. and G., P. (2010). Analyzing the Amazon Mechanical Turk market-place. *XRDS: Crossroads, The ACM Magazine for Students*, 17(2):16–21.
- Israel, G. D. (1992). Determining Sample Size 1. *Florida Cooperative Extension Service, University of Florida*, PEOD-6.
- Kiefer, P., Giannopoulos, I., Duchowski, A., and Raubal, M. (2016). Measuring Cognitive Load for Map Tasks Through Pupil Diameter. pages 323–337. Springer, Cham.
- Kitchin, R. and Tate, N. J. (2000). *Conducting Research into Human Geography*. Prentice Hall.
- Kostas (2016). Using Crowdsourcing and Machine Learning to locate swimming pools in Australia · Tomnod. Date Accessed: 2017-05-04 URL: blog.tomnod.com/crowd-and-machine-combo.
- LaMorte, W. W. (2017). Mann Whitney U Test (Wilcoxon Rank Sum Test). Date Accessed: 2017-05-12 URL: sphweb.bumc.bu.edu/otlt/mphs-modules/bs/bs704_nonparametric/BS704_Nonparametric4.html.
- Leppink, J., Paas, F., Van Gog, T., Van Der Vleuten, C. P. M., and Van Merriënboer, J. J. G. (2014). Effects of pairs of problems and examples on task performance and different types of cognitive load. *Learning and Instruction*, 30:32–42.
- Lieu, D. (2017). End-to-End Satellite Imagery Analysis — Development Seed. Date Accessed: 2017-06-03 URL: developmentseed.org/blog/2017/04/21/end-to-end-ml/.

BIBLIOGRAPHY

- Lund Research Ltd (2013a). Hypothesis Testing - Significance levels and rejecting or accepting the null hypothesis. Date Accessed: 2017-04-20 URL: statistics.laerd.com/statistical-guides/hypothesis-testing-3.php.
- Lund Research Ltd (2013b). Mann-Whitney U Test in SPSS Statistics | Setup, Procedure & Interpretation | Laerd Statistics. Date Accessed: 2017-05-12 URL: statistics.laerd.com/spss-tutorials/mann-whitney-u-test-using-spss-statistics.php.
- Lund Research Ltd (2013c). One-way ANOVA - Its preference to multiple t-tests and the assumptions needed to run this test | Laerd Statistics. Date Accessed: 2017-04-25 URL: statistics.laerd.com/statistical-guides/one-way-anova-statistical-guide-2.php.
- Mandler, G. (2013). The Limit of Mental Structures, The Journal of General Psychology. pages 243–250.
- MedCalc Software bvba (2017). Skewness and Kurtosis. Date Accessed: 2017-04-23 URL: medcalc.org/manual/skewnesskurtosis.php.
- Meier, P. (2013a). Digital Humanitarian Response: Moving from Crowd-sourcing to Microtasking | iRevolutions. Date Accessed: 2017-05-04 URL: irevolutions.org/2013/01/20/digital-humanitarian-micro-tasking/.
- Meier, P. (2013b). Handbook of Human Computation. In *Handbook of Human Computation*. Springer New York, New York, NY.
- Meier, P. (2014). Typhoon | iRevolutions. Date Accessed: 2017-05-07 URL: irevolutions.org/tag/typhoon/.
- Michelucci, P. and Dickinson, J. L. (2016). The power of crowds. *Science*, 351(6268):32–33.
- Morschheuser, B., Hamari, J., and Koivisto, J. (2016). Gamification in Crowdsourcing: A Review. In *2016 49th Hawaii International Conference on System Sciences (HICSS)*, pages 4375–4384. IEEE.
- Motulsky, H. (2013). Intuitive Biostatistics: A Nonmathematical Guide to Statistical Thinking, 3rd edition. In *Intuitive Biostatistics*, chapter 24.
- Nakhuda, D. (2016). Announcing General Availability of Mighty AI's Intelligent Crowdsourcing Platform | Mighty AI. Date Accessed: 2017-05-23 URL: mty.ai/blog/general-availability-intelligent-crowdsourcing-platform/.
- Nikki (2016). Finding Swimming Pools in Australia using Deep Learning. Date Accessed: 2017-05-04 URL: blog.tonmod.com/finding-pools-with-deep-learning.
- Oppenheimer, D. (2017). Machine Learning with Humans in the Loop - Algorithmia. Date Accessed: 2017-05-14 URL: blog.algorithmia.com/machine-learning-with-human-in-the-loop/.

-
- Ørstavik, M. (2017). AirNet – A deep learning approach to extracting building information from remotely sensed images.
- Osborne, J. W. (2010). Improving your data transformations: Applying the Box-Cox transformation. *Practical Assessment, Research & Evaluation*, 15(12).
- Palen, L., Soden, R., Anderson, T. J., and Barrenechea, M. (2015). Success & Scale in a Data-Producing Organization. *Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems - CHI '15*, pages 4113–4122.
- Pearson, A., Sözcükler, A., düzeltmeli Kolmogorov-Smirnov, L., Pearson ve Jarquata-Bera testleri Derya ÖZTUNA Atilla Halil ELHAN Ersöz TÜCCAR, A., and Öz-tuna, D. (2006). Investigation of Four Different Normality Tests in Terms of Type 1 Error Rate and Power under Different Distributions. *Turk J Med Sci*, 36(3):171–176.
- Quinn, A. J. and Bederson, B. B. (2011). Human computation. *Proceedings of the 2011 annual conference on Human factors in computing systems - CHI '11*, pages 1403–1412.
- Ross, C. and Jamilly, M. (2016). Better Together how Computers can be Designed to Augment Human Ability. *ITNOW*, 58(2):44–47.
- Salk, C., Sturn, T., See, L., and Fritz, S. (2016). Local Knowledge and Professional Background Have a Minimal Impact on Volunteer Citizen Science Performance in a Land-Cover Classification Task. *Remote Sensing*, 8(10):774.
- Sambale, M. (2016). OpenStreetMap | manings sin dagbok | Building tools for LABuildings Import. Date Accessed: 2017-05-31 URL: openstreetmap.org/user/manings/diary/38969.
- Sarasua, C., Simperl, E., and Noy, N. F. (2012). Crowdsourcing Ontology Alignment with Microtasks. pages 525–541.
- Schade, A. (2015). Pilot Testing: Getting It Right (Before) the First Time. Date Accessed: 2017-04-19 URL: nngroup.com/articles/pilot-testing/.
- Schulze, T., Krug, S., and Schader, M. (2012). Workers' Task Choice in Crowdsourcing and Human Computation Markets. *ICIS 2012 Proceedings*.
- See, L., Comber, A., Salk, C., Fritz, S., van der Velde, M., Perger, C., Schill, C., McCallum, I., Kraxner, F., and Obersteiner, M. (2013). Comparing the quality of crowdsourced data contributed by expert and non-experts. *PloS one*, 8(7).
- Smith, S. (2013). Determining Sample Size: How to Ensure You Get the Correct Sample Size | Qualtrics. Date Accessed: 2017-04-19 URL: qualtrics.com/blog/determining-sample-size/.
- Stanford University (2017). Machine Learning | Coursera. Date Accessed: 2017-05-14 URL: www.coursera.org/learn/machine-learning.

BIBLIOGRAPHY

- The Pennsylvania State University (2017). 7.5 - Power and Sample Size Determination for Testing a Population Mean | STAT 500.
- The Scipy community (2017). `scipy.stats.anderson` — SciPy v0.19.0 Reference Guide. Date Accessed: 2017-04-21 URL: docs.scipy.org/doc/scipy/reference/generated/scipy.stats.anderson.html.
- von Ahn, L. (2008). Human Computation. In *2008 IEEE 24th International Conference on Data Engineering*, pages 1–2. IEEE.
- Walpole, R. E., Myers, R. H., Myers, S. L., and Ye, K. (2012). *Probability & Statistics*. Pearson Education, ninth edition.
- Wang, X., Goh, D. H.-L., Lim, E.-P., Wei Liang Vu, A., and Chua, A. Y. K. (2017). Examining the Effectiveness of Gamification in Human Computation. *International Journal of Human-Computer Interaction*, pages 1–9.
- Webster, E. (2016). Using Crowdsourcing for Complex Data Problems: It Can & Should Be Done! | Mighty AI. Date Accessed: 2017-05-23 URL: mty.ai/blog/using-crowdsourcing-for-complex-data-problems-it-can-should-be-done/.
- Yang, J., Redi, J., Demartini, G., and Bozzon, A. (2016). Modeling Task Complexity in Crowdsourcing. *ResearchGate*.
- Yap, B. W. and Sim, C. H. (2011). Journal of Statistical Computation and Simulation Comparisons of various types of normality tests Comparisons of various types of normality tests. *Journal of Statistical Computation and Simulation*, 81(12):2141–2155.